

CS 314 – Operating Systems

Caveat

- Expect to be busy
- Assignments take time
 - Plan accordingly
 - E.g. if 3 weeks are given for an assignment, do not wait until the last couple of days to start
- Keep up
 - Material is not difficult, but can be plentiful
- Know the syllabus
- Interact
 - If you have questions, ask

“Study the past if you would define the future”
-Confucius

- Former 314 students' thoughts about assignments
 - “Start early on the labs and homework”
 - “Get a partner for labs, do them incrementally”
 - “the programs... might make your brain melt”
 - “give yourself enough time to do the labs and ask a lot of questions if you are stuck”
 - “The programs are not as easy as one would hope”
 - “The labs take a lot of time.”
 - “Take the time to understand the assignments, then they will be much easier”

“Sweet is the memory of past troubles”
-Faulkner

- Former 314 students' thoughts about the environment:
 - “Know everything”
 - “get to know working in a UNIX environment early in the semester”
 - “Invest time in learning bash and C before any labs are assigned. If a lab is due in two weeks, work all of those two weeks on it.”

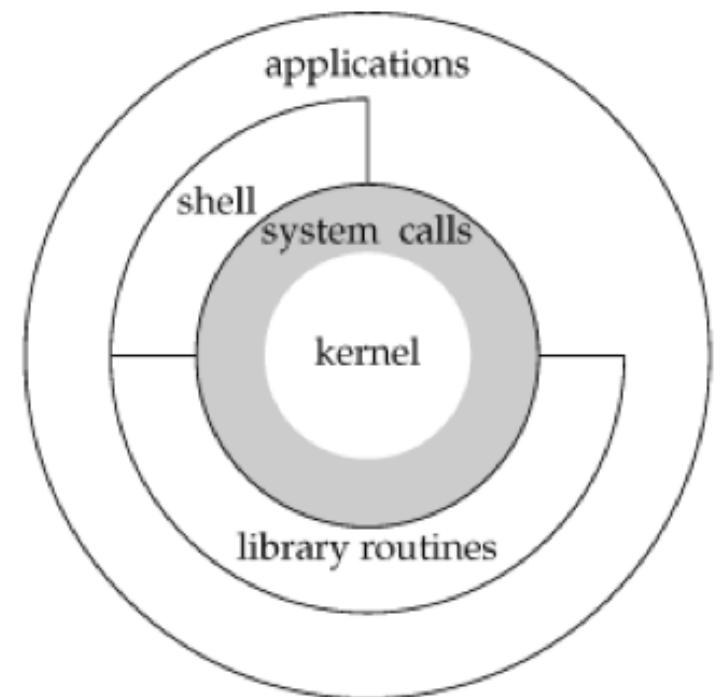
“There is no present or future – only the past, happening over and over again – now”

-Eugene O'Neill

- Former 314 students' thoughts about stuff in general:
 - “Go talk to Professor Crk, no one does and he wants to help”
 - “Go to office hours”
 - “CS 314 is a 2 semester course, they just don't mention that in the catalog description”
 - “The internet is a useful tool, but don't use it as your only tool”
 - “Read the book. Come to class. DON'T PROCRASTINATE”
 - “Don't put off assignments in this class like you might do in others”
 - “Look through the slides to determine what to read.”

The Simple Big Picture

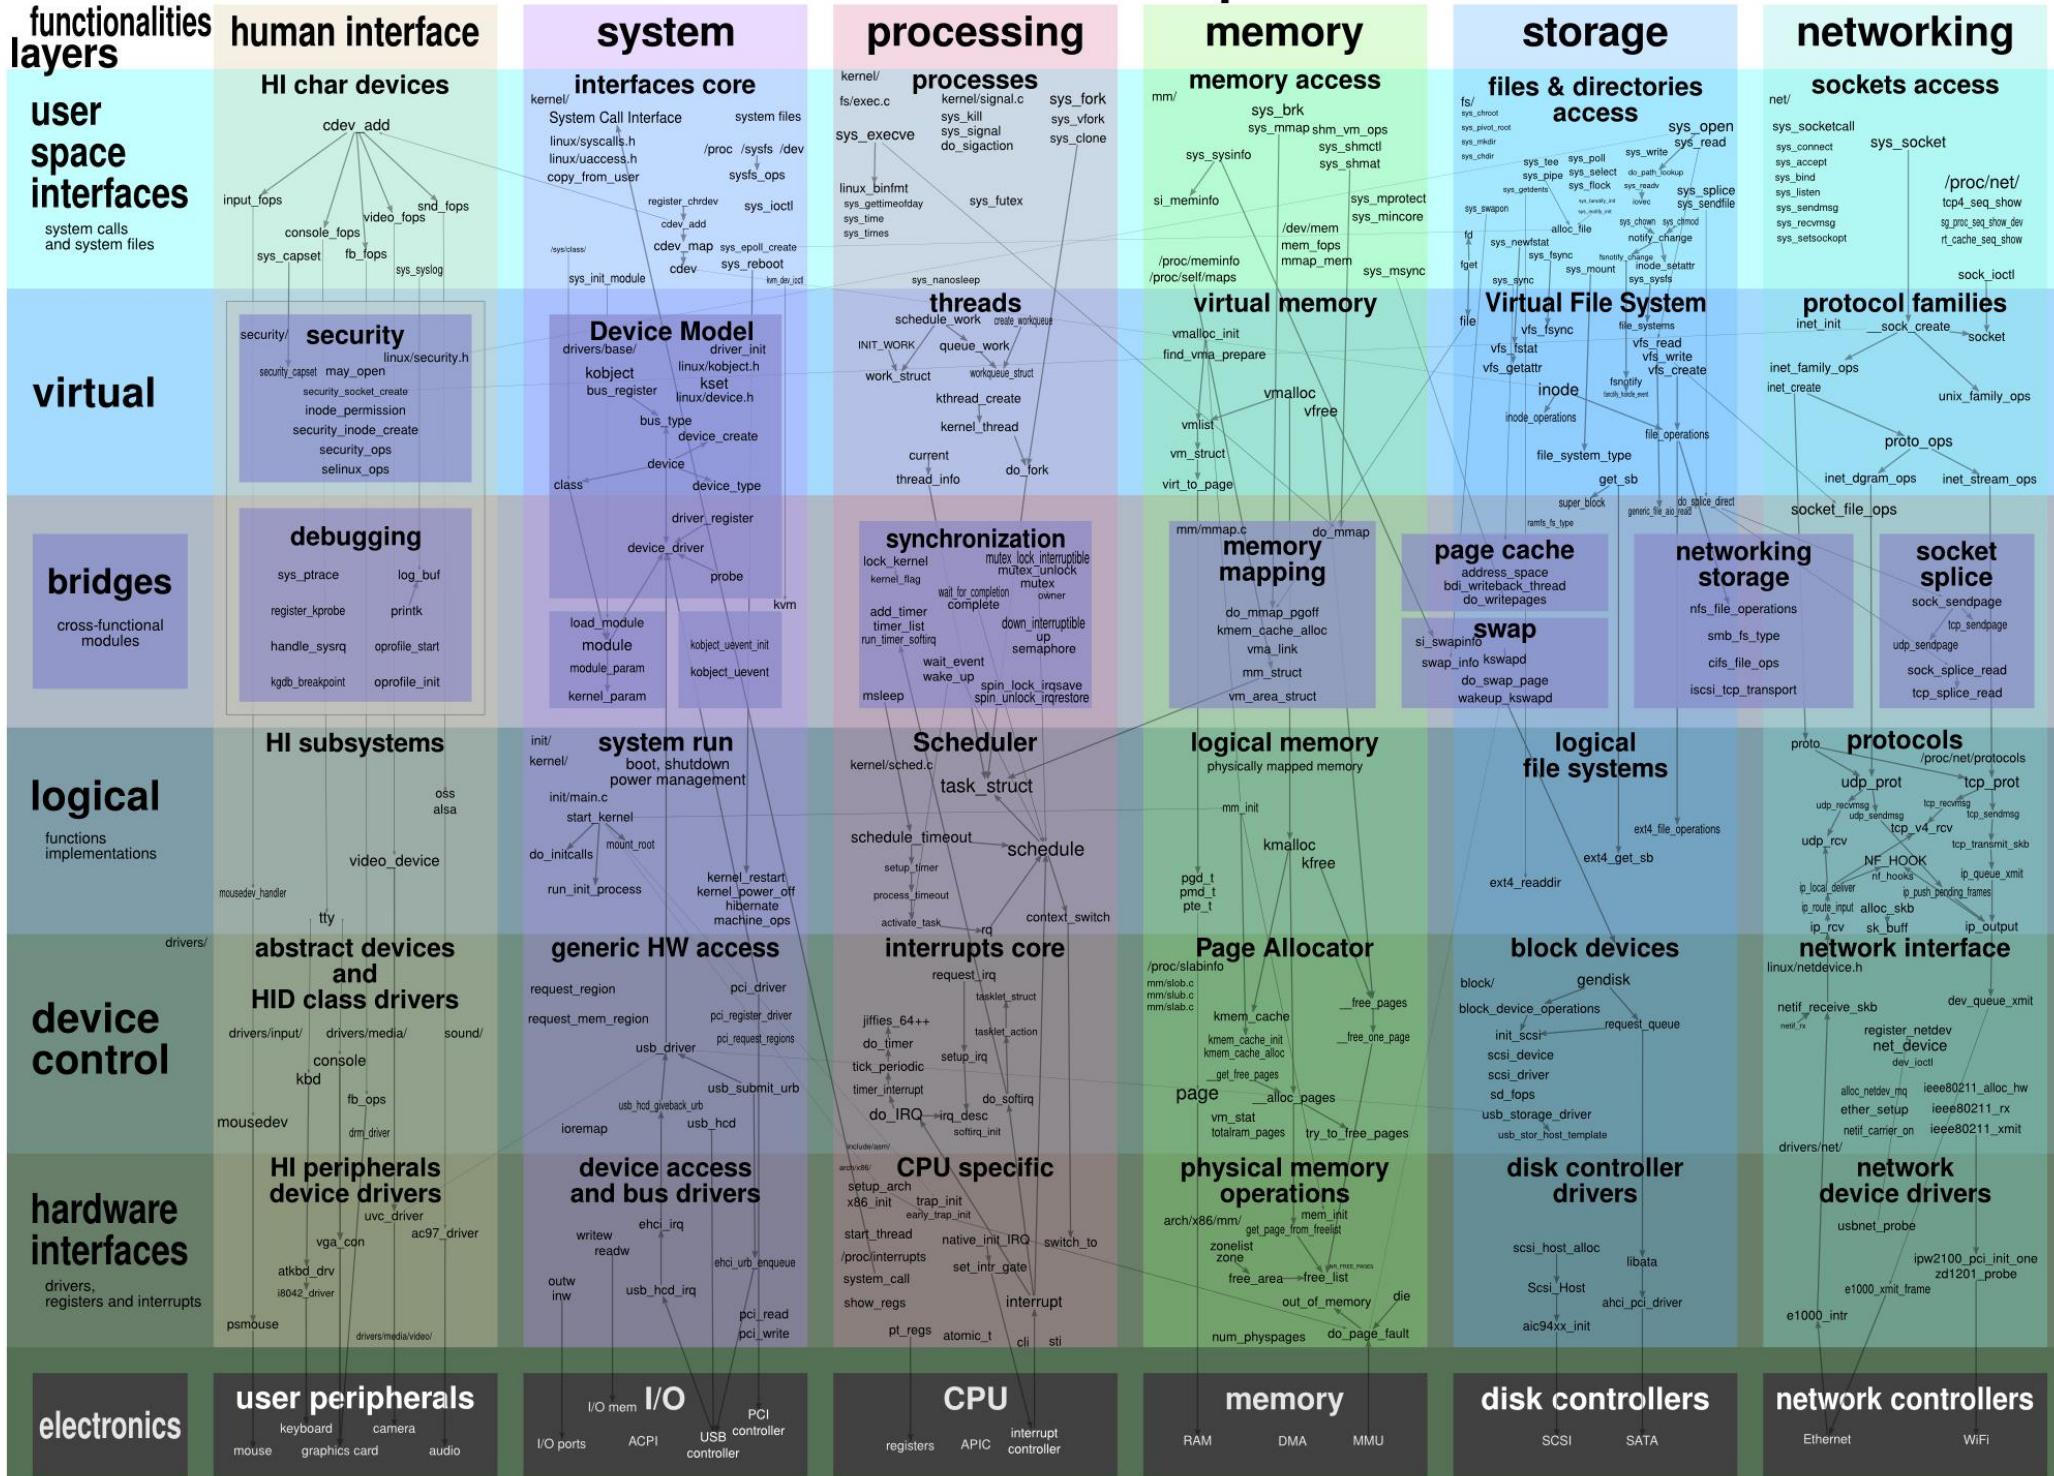
- Kernel: resource management
- System calls: how we ask the kernel action
- Shell/UI: how we interact with the system
- Library routines: how we build functionality around system calls
- Applications: programs whose timely and efficient execution users care about



OS – What are we talking about?

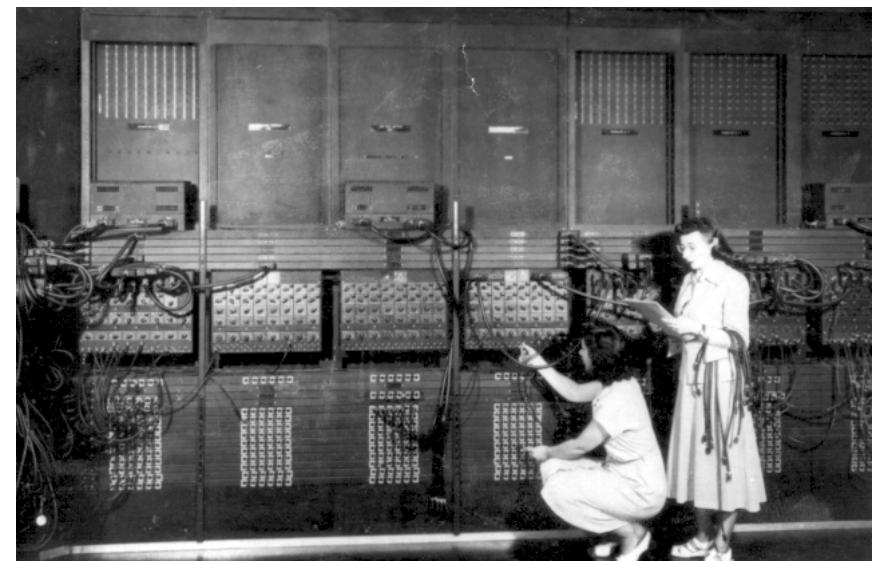
- OS is an extended machine
 - architecture is primitive and/or awkward
 - OS presents a virtual machine that should be easier to program
- OS does resource management
 - How do we provide the means for multiple programs to run?
 - How do we manage and protect memory/devices?
 - How do we multiplex resources in time and space?

Linux kernel map



How did we get here? (A general history)

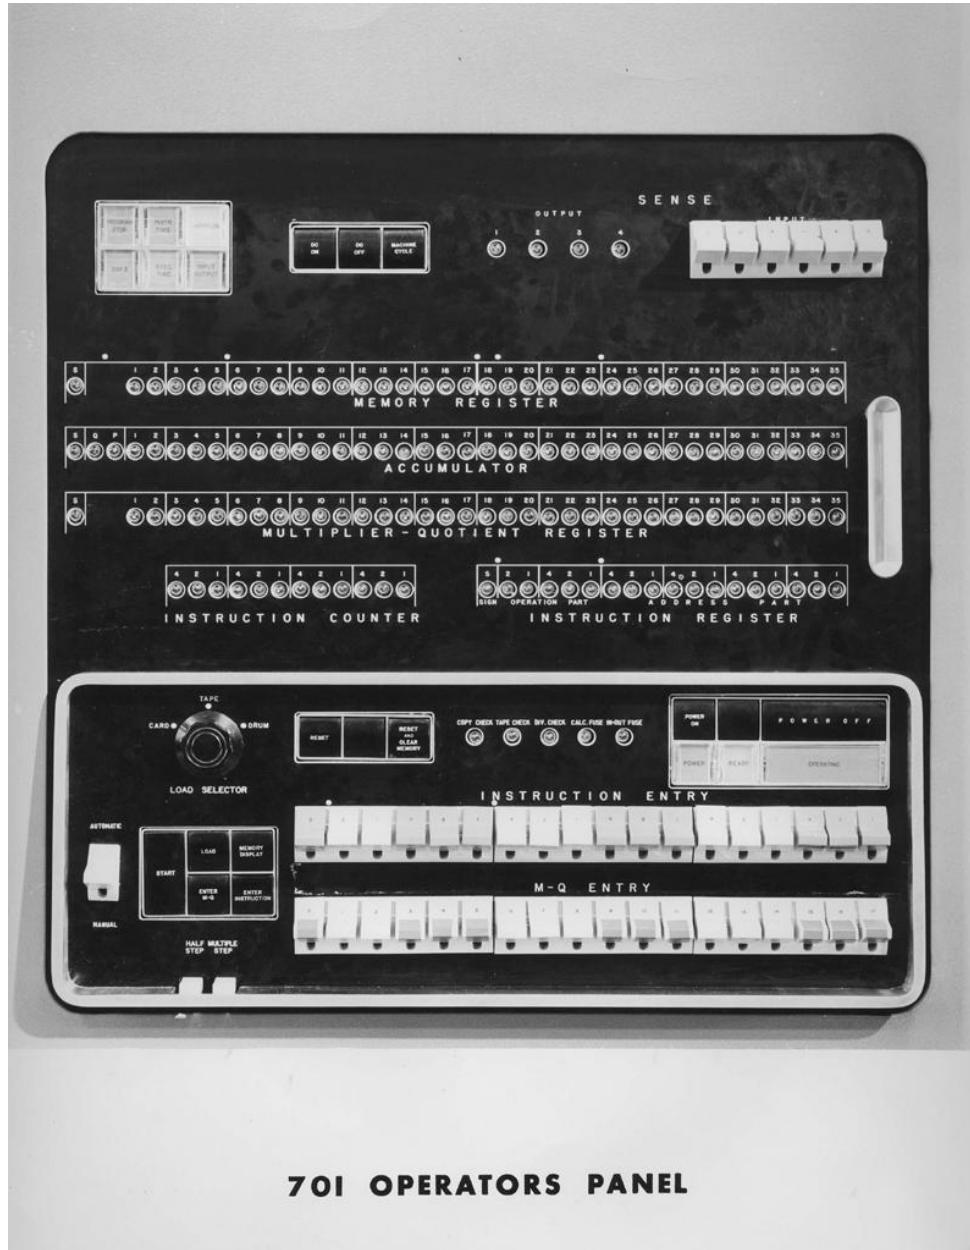
- All advances drive towards greater automation
- 1945-1955 (vacuum tubes)
 - No OS, no programming languages
 - Everything is a manual effort
 - Absolute machine language or directly wiring circuits
 - Programmers reserved time to insert a plugboard
 - Punch cards in 1950s



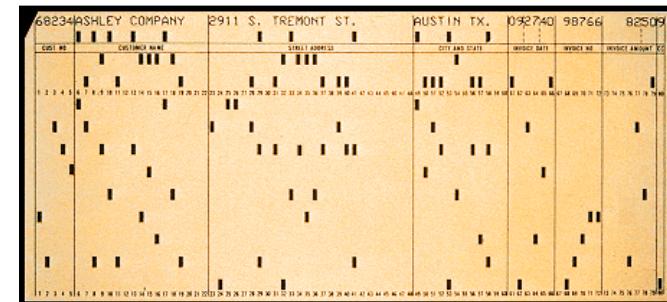
What the 50's looked like (IBM 701)



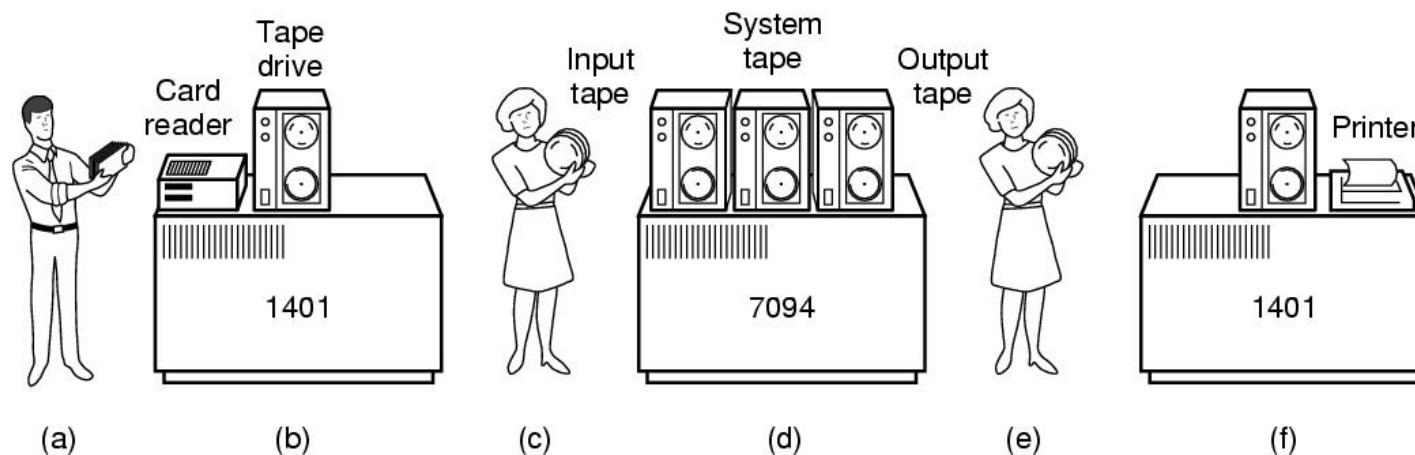
What the 50's looked like (IBM 701)



and so on...



- 1955-1965 (transistors, batch systems)
 - first clear separation between designers, operators, programmers, etc.
 - programmers hand off stacks of punch cards to operators, do something else while waiting for output
 - operators carried tapes containing batches of jobs from card readers to the computing machine and again out to the output printer



What the 60's looked like (IBM 7094)



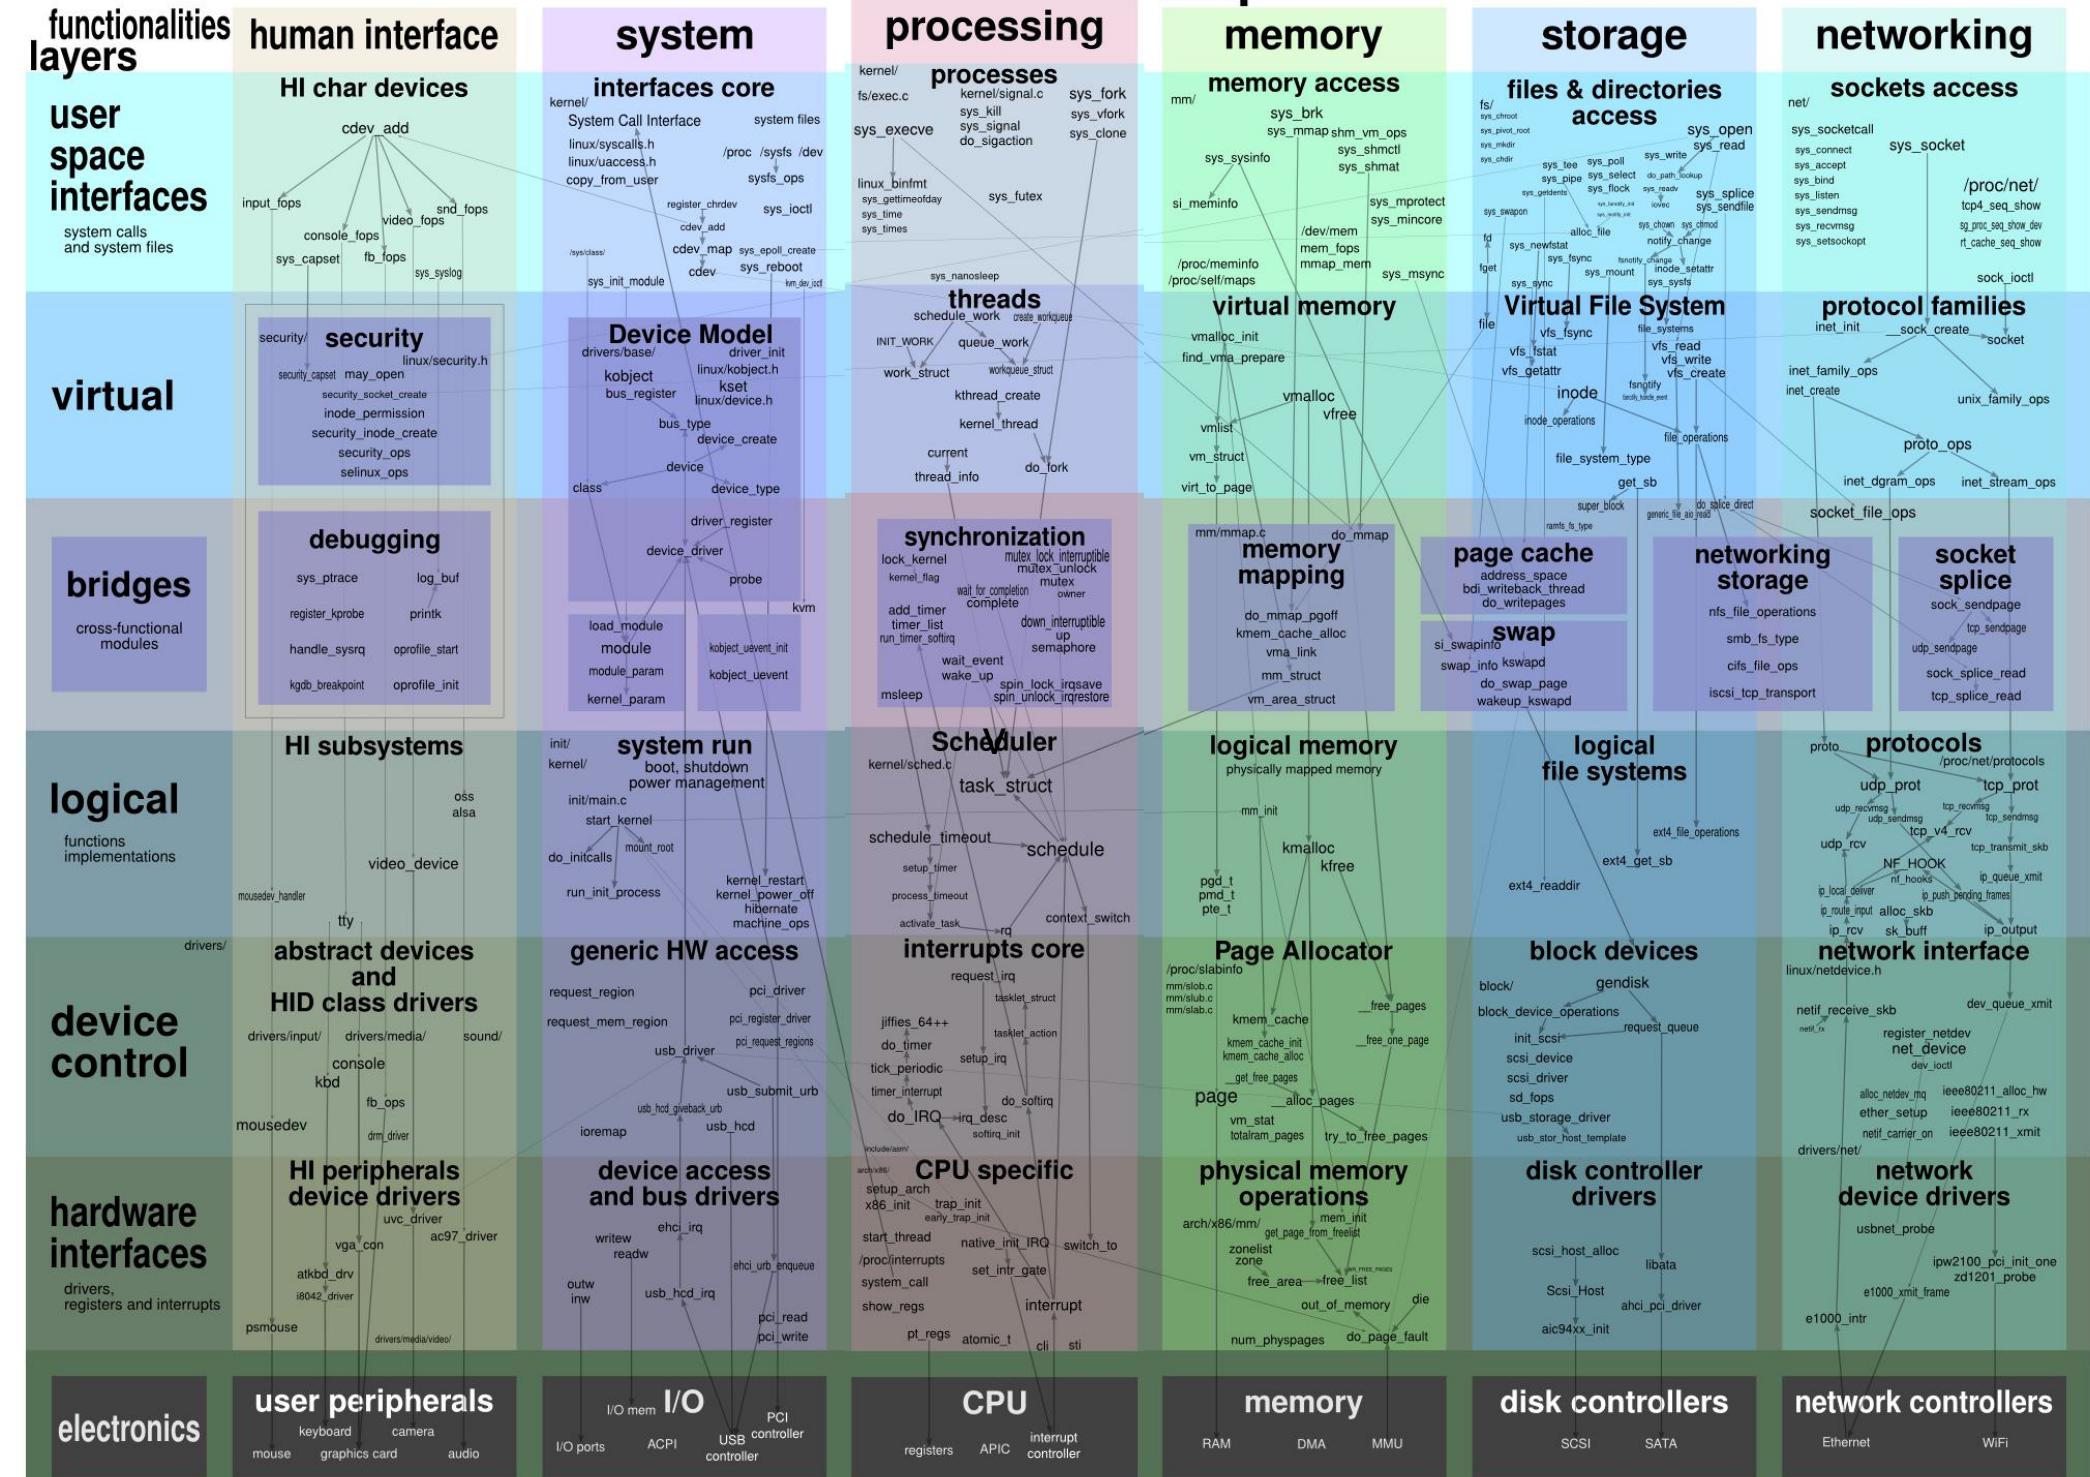
What the 60's looked like (IBM 7094)



and so on...

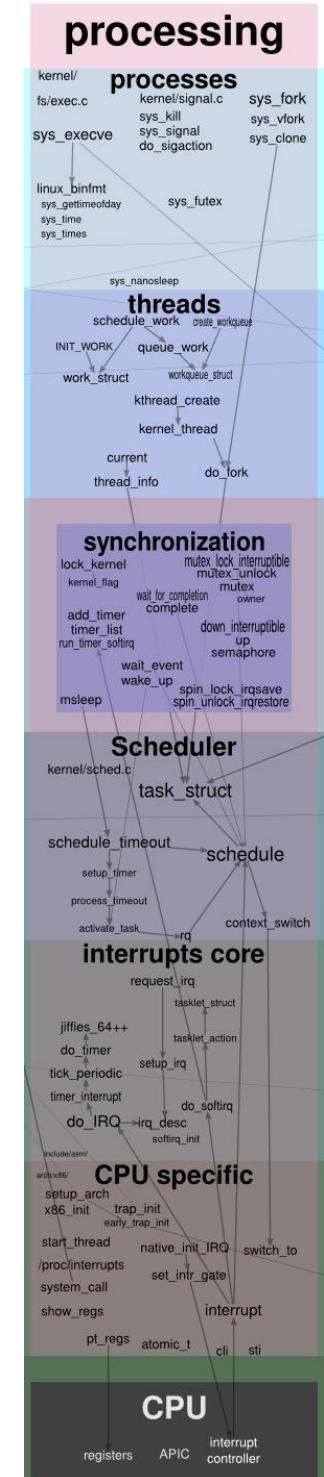
- 1965-1980 (Multiprogramming)
 - While a job is waiting for some I/O operation to complete, allow another job to use the CPU
 - Timesharing – allowing many simultaneous users
- 1980-Present
 - Personal computing
 - DOS, Windows, Linux, Android, OSX, etc.
- In short
 - all modern operating systems designs come from early ones
 - all solve essentially the same resource management problems: CPU sharing, memory partitioning/sharing, I/O device sharing, interfacing with users and programs

Linux kernel map



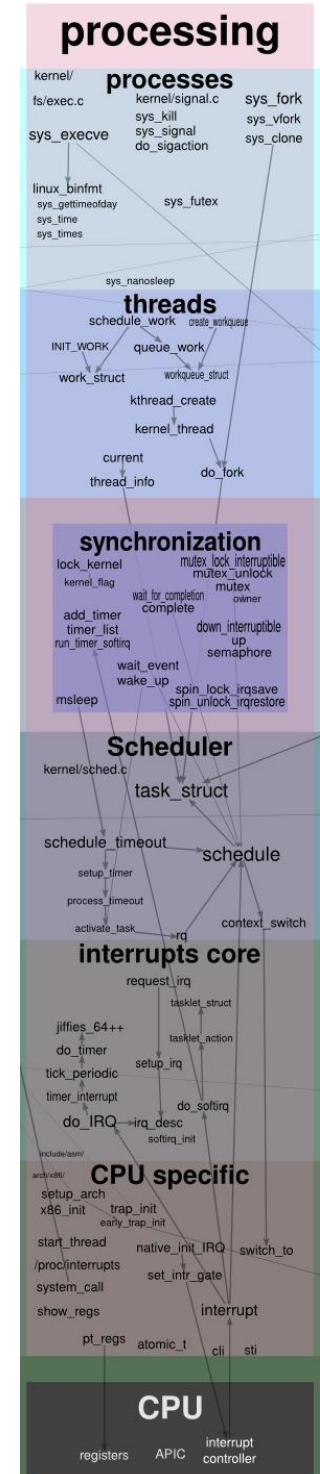
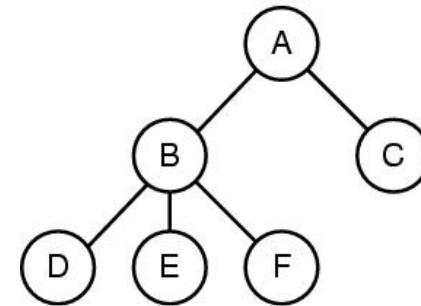
CPU Review

- Fetch/execute instructions
 - has unique instruction set
- General and special registers
 - Program counter (PC): keep track of current inst.
 - Stack pointer (SP): currently executing procedure
 - Program status word (PSW): what is the program allowed to do
 - user/kernel mode, CPU priority, etc...



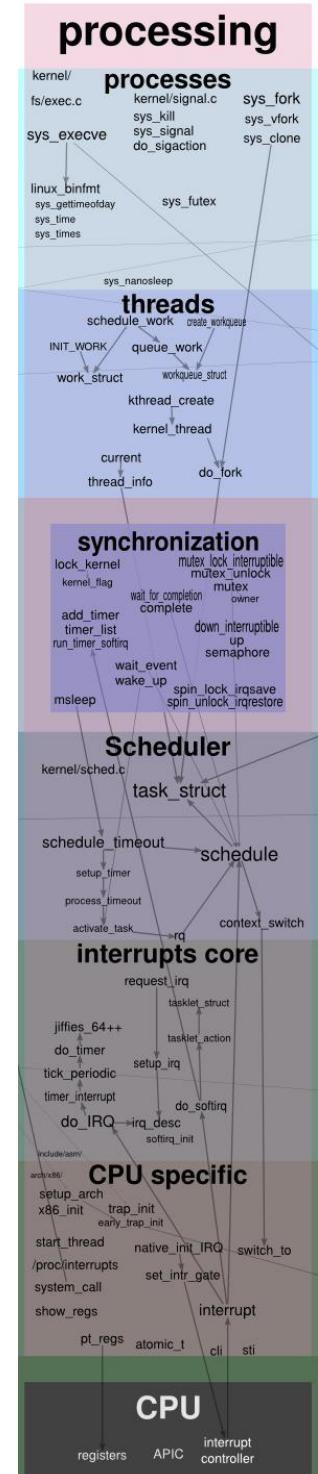
Processes in General

- Processes are programs in execution
- From the OS perspective
 - process = container defining an address space which contains:
 - executable code
 - program data
 - program stack
- OS manages processes
- Processes can reproduce
- Processes can communicate



Processes in General

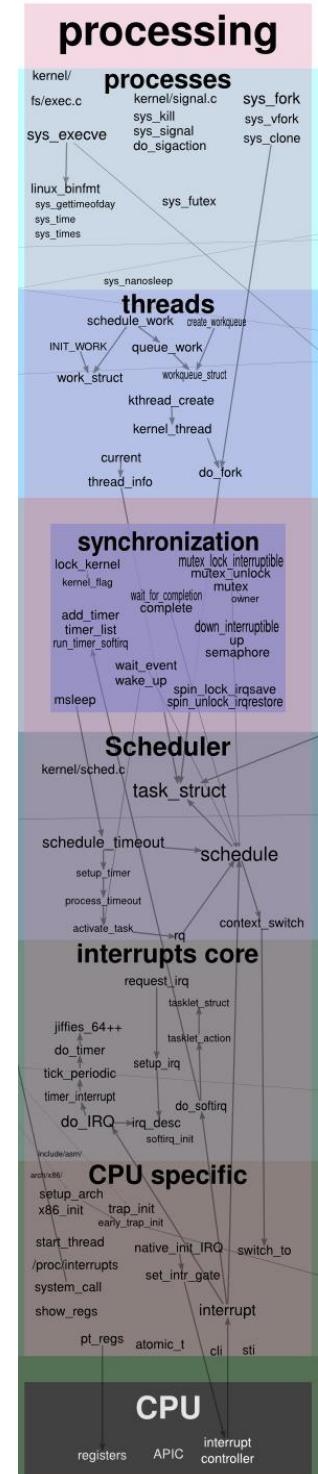
- Multiuser systems allow several processes to be active concurrently
- Systems allowing concurrent active processes are multiprogramming systems
- scheduler decides which process is allowed to progress
 - OS may be nonpreemptive
 - scheduler invoked only when CPU is voluntarily relinquished
 - or preemptive
 - scheduler invoked by OS periodically, a necessity for multiuser systems



Processes in General

- (in UNIX) Processes are under the illusion that they are the only process on the machine with exclusive access to the kernel
 - kernel itself is not a process, but a process manager
- Threads allow one process to have multiple PCs and SPs, or multiple processes to work within the same address space...
- Creating a new process in Unix is simple

```
int pid = fork();
if (pid == 0) {
    exec("foo");
} else {
    waitpid(pid, &status, options);
};
```



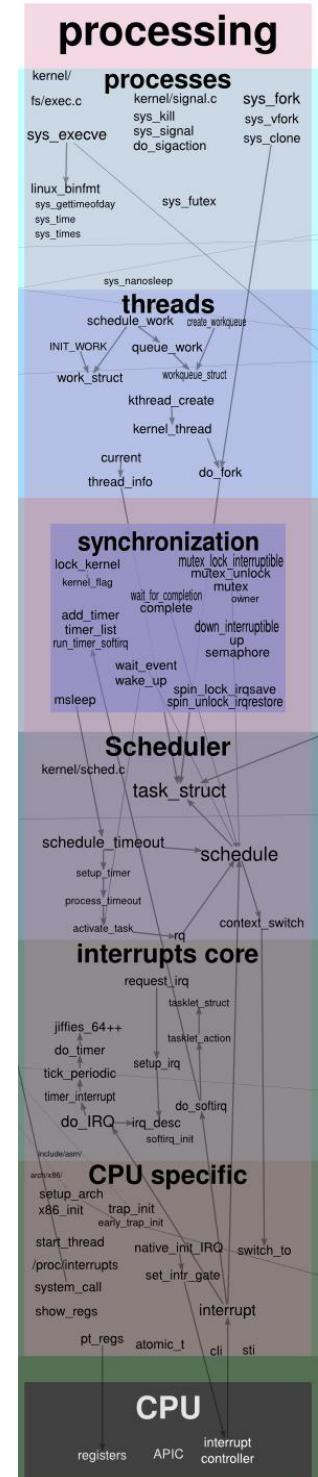
Processes in General

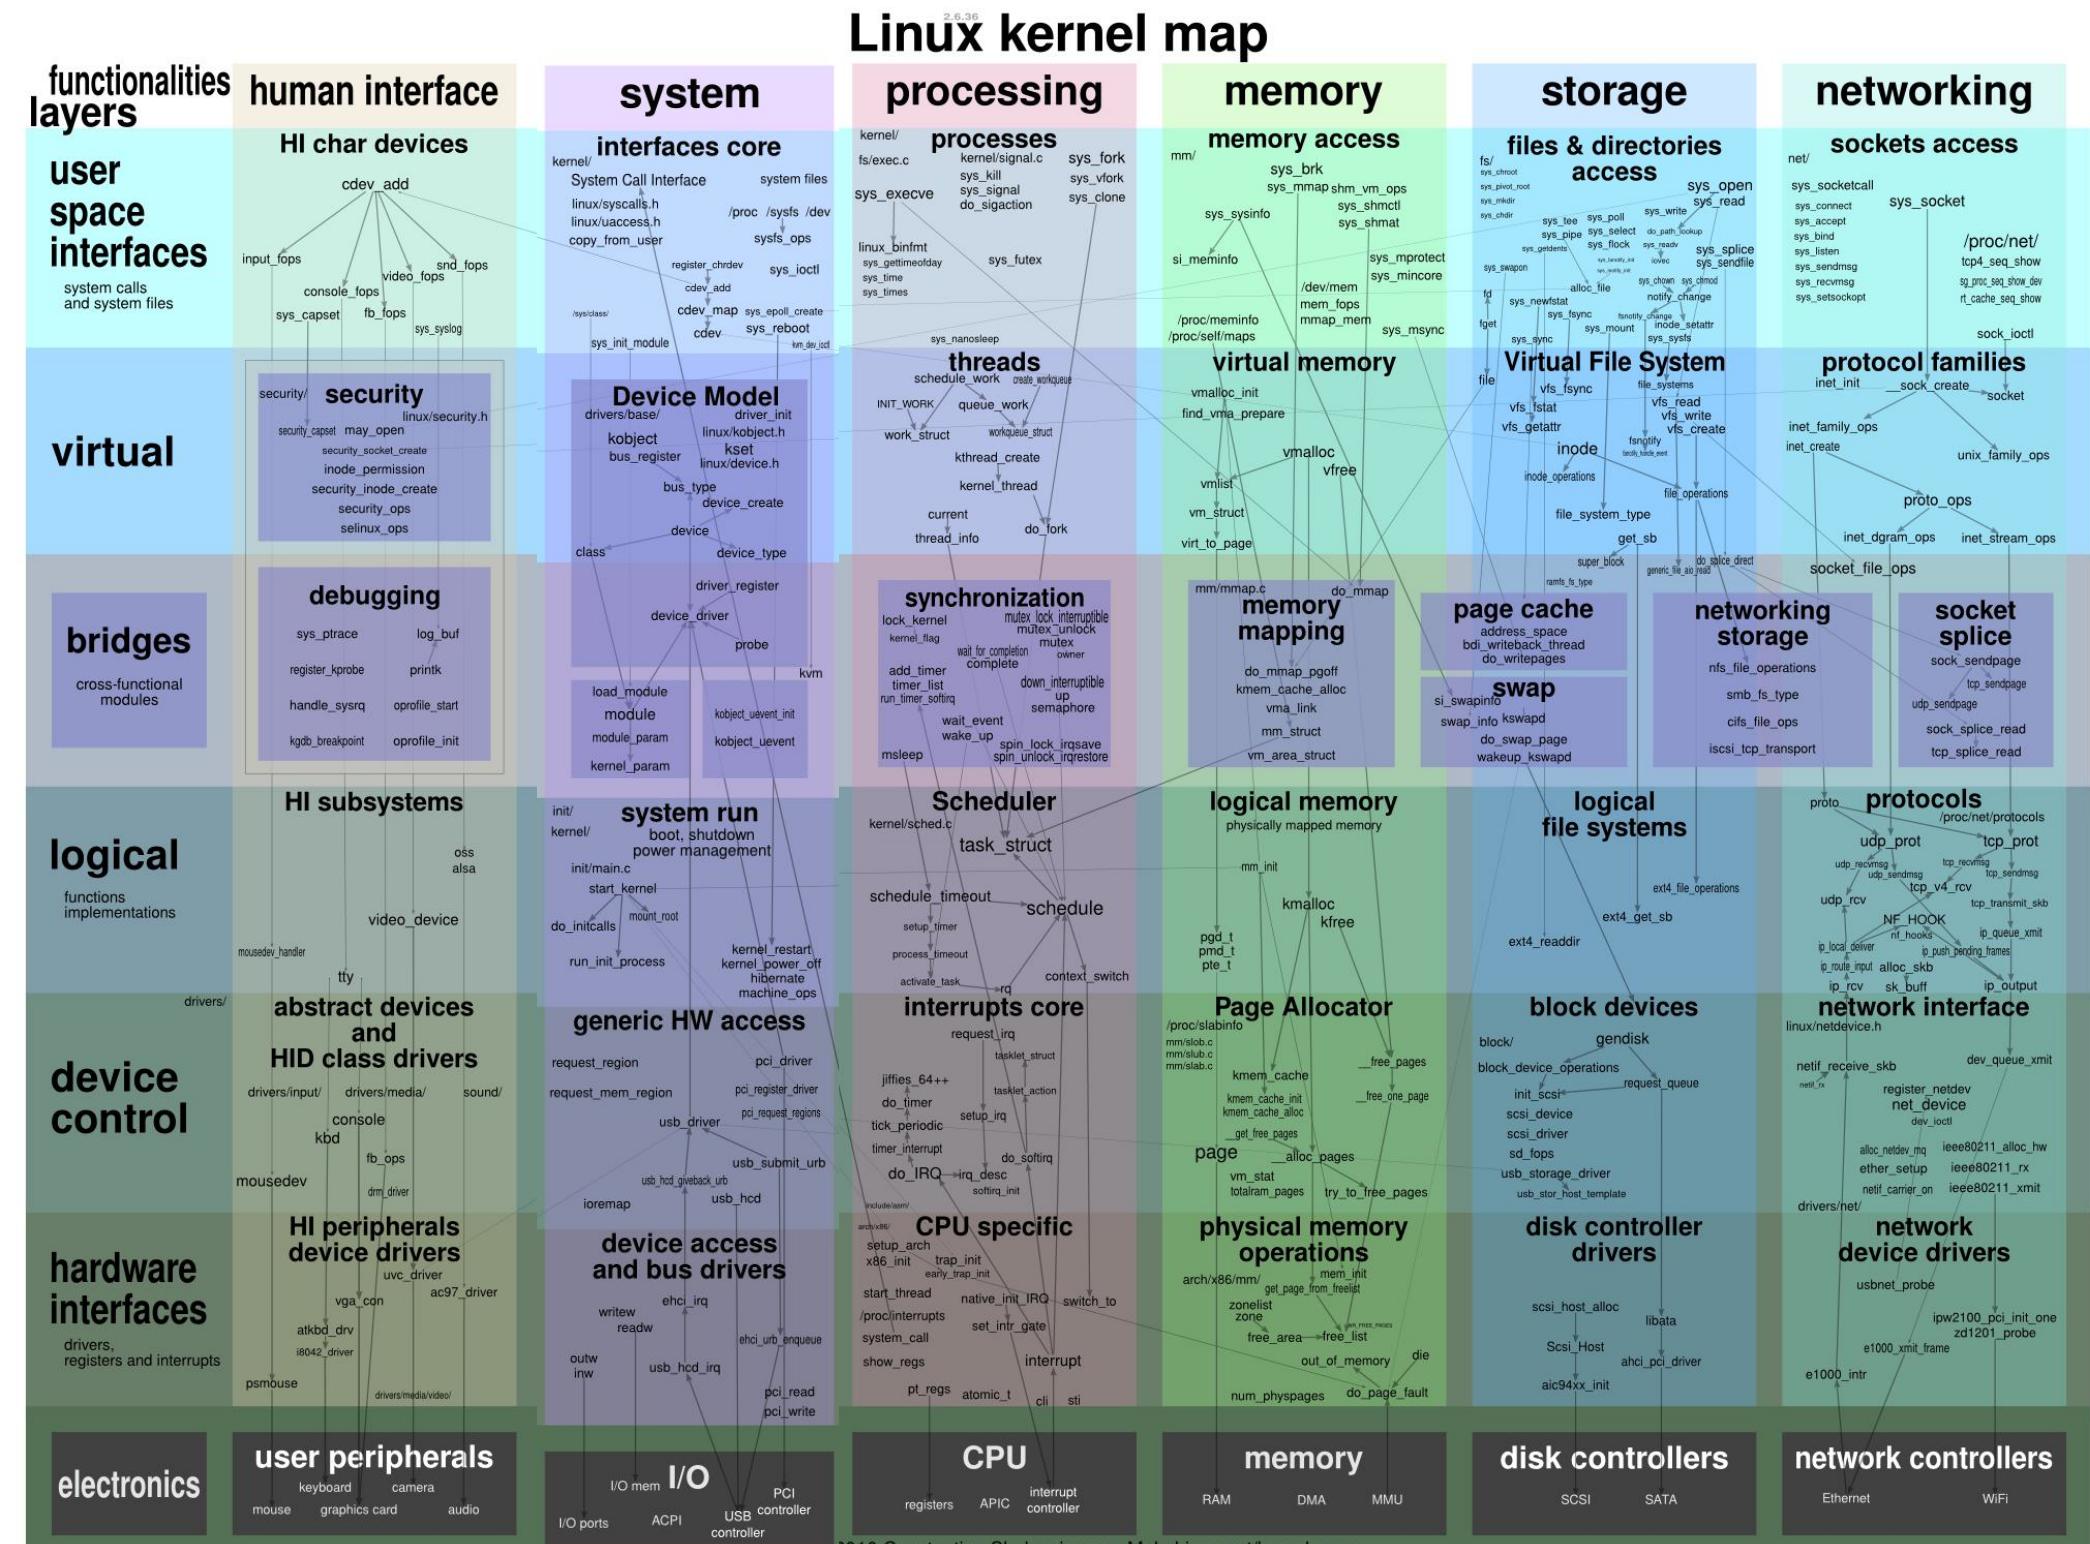
- In windows, it's arguably not:

```
BOOL CreateProcess(
    LPCTSTR lpApplicationName,
    LPTSTR lpCommandLine,
    LPSECURITY_ATTRIBUTES lpProcessAttributes,
    LPSECURITY_ATTRIBUTES lpThreadAttributes,
    BOOL bInheritHandles,
    DWORD dwCreationFlags,
    LPVOID lpEnvironment,
    LPCTSTR lpCurrentDirectory,
    LPSTARTUPINFO lpStartupInfo,
    LPPROCESS_INFORMATION lpProcessInformation
);

WaitForSingleObject(lpProcessInformation->hProcess,
    INFINITE);
```

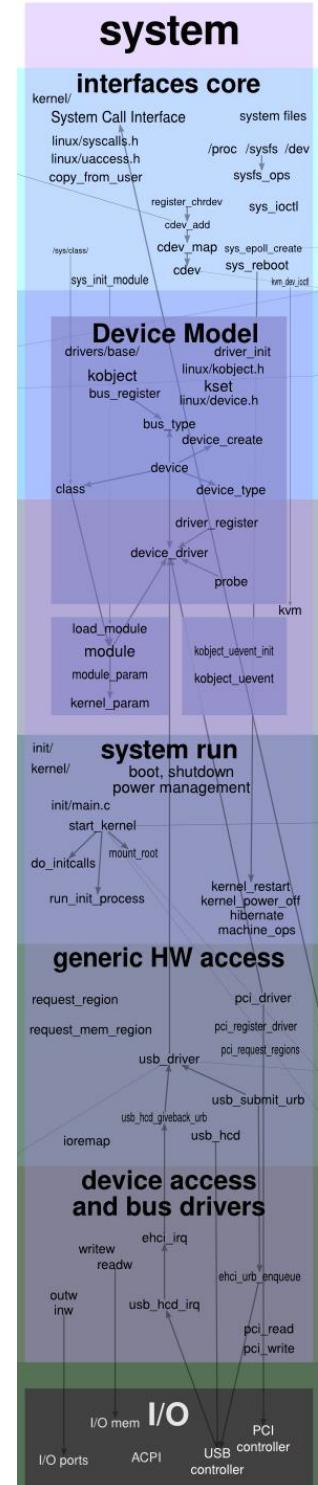
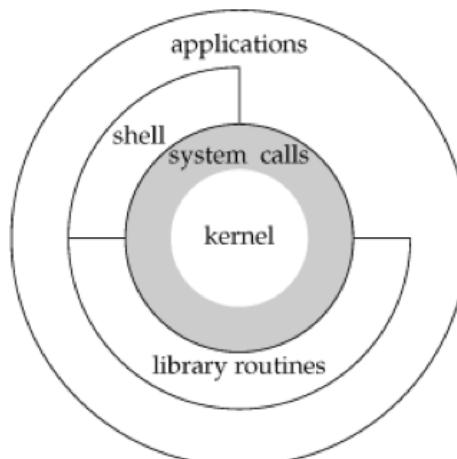
- So... we'll stick with Unix-like systems
 - Note: this reason alone would be insufficient, but consider it illustrative of the situation

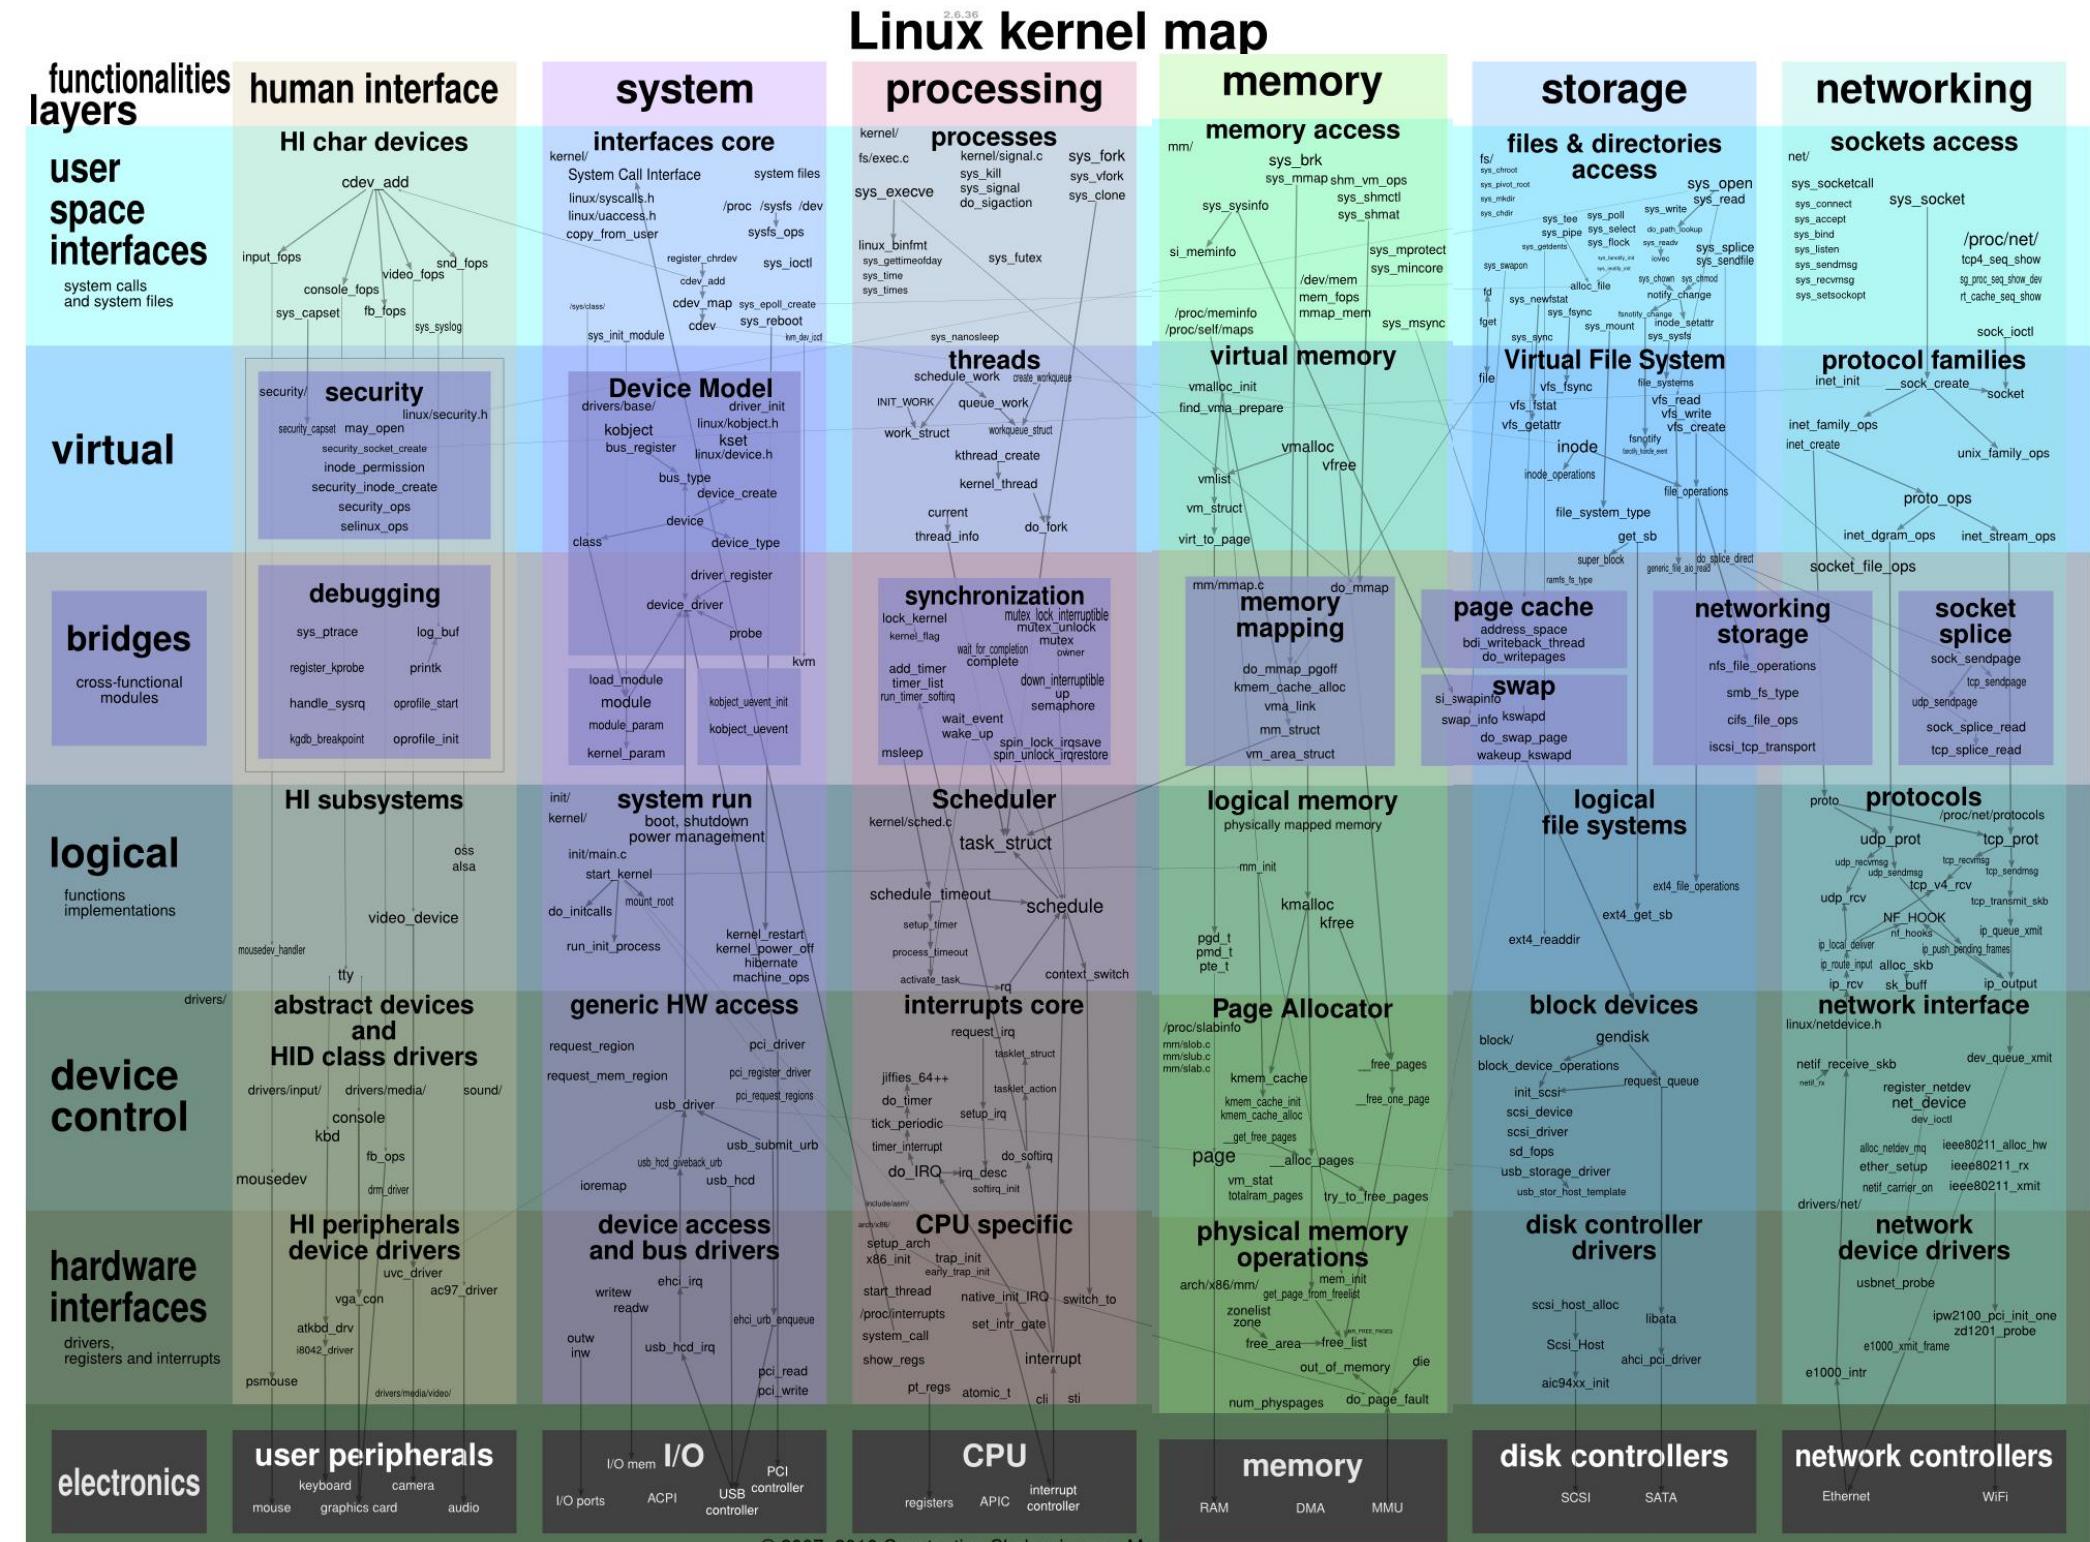




System Calls

- Processes make system calls to request OS services
 - privilege mode is changed from user mode to kernel mode
 - the intended kernel procedure is executed
 - on completion, kernel procedure returns HW to user mode and process continues where it left off

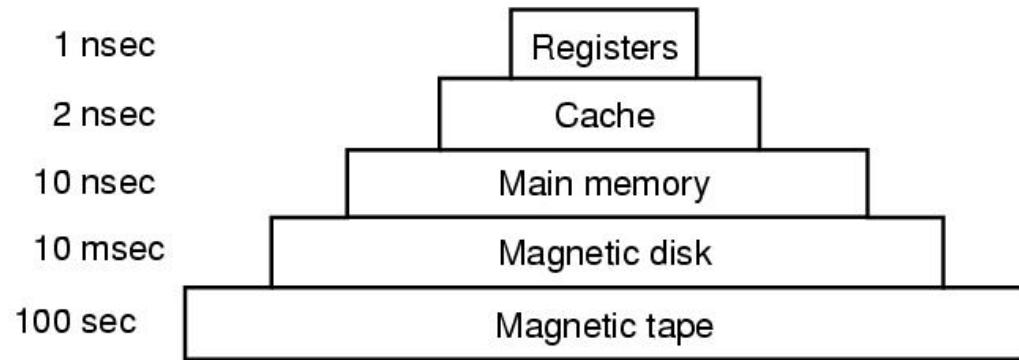




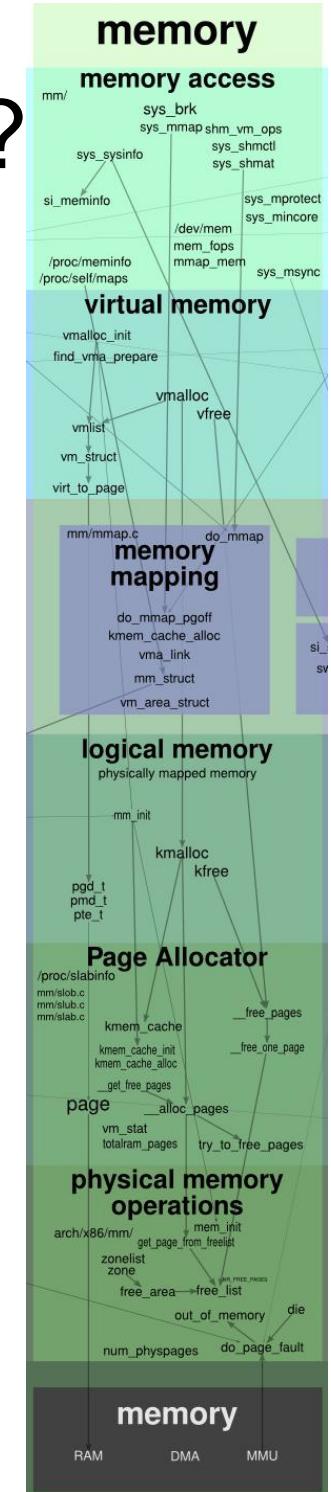
Remember the Memory Hierarchy?

- Capacities and access times are approximate

Typical access time

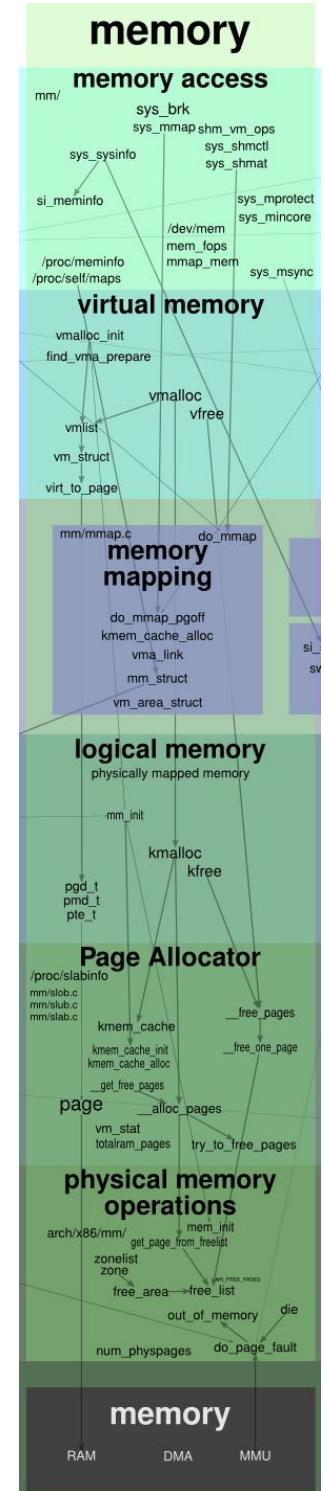


Typical capacity



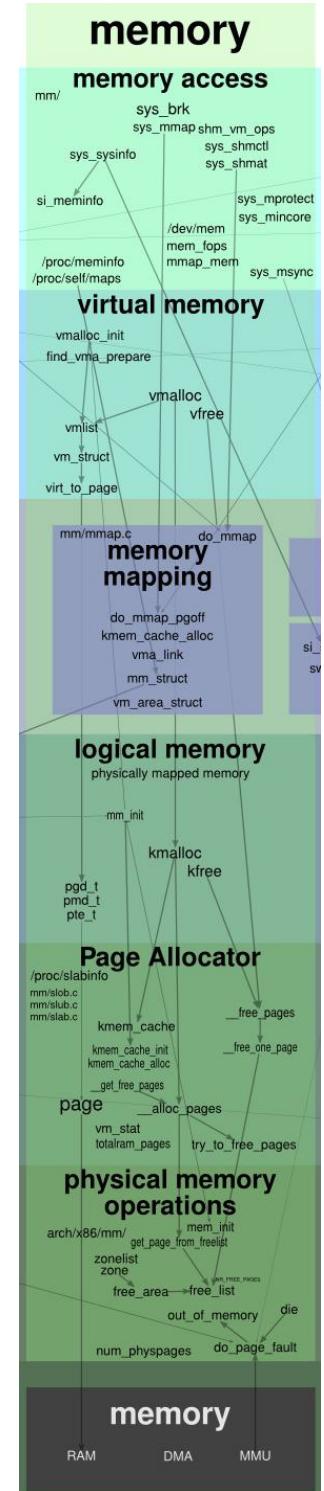
So what?

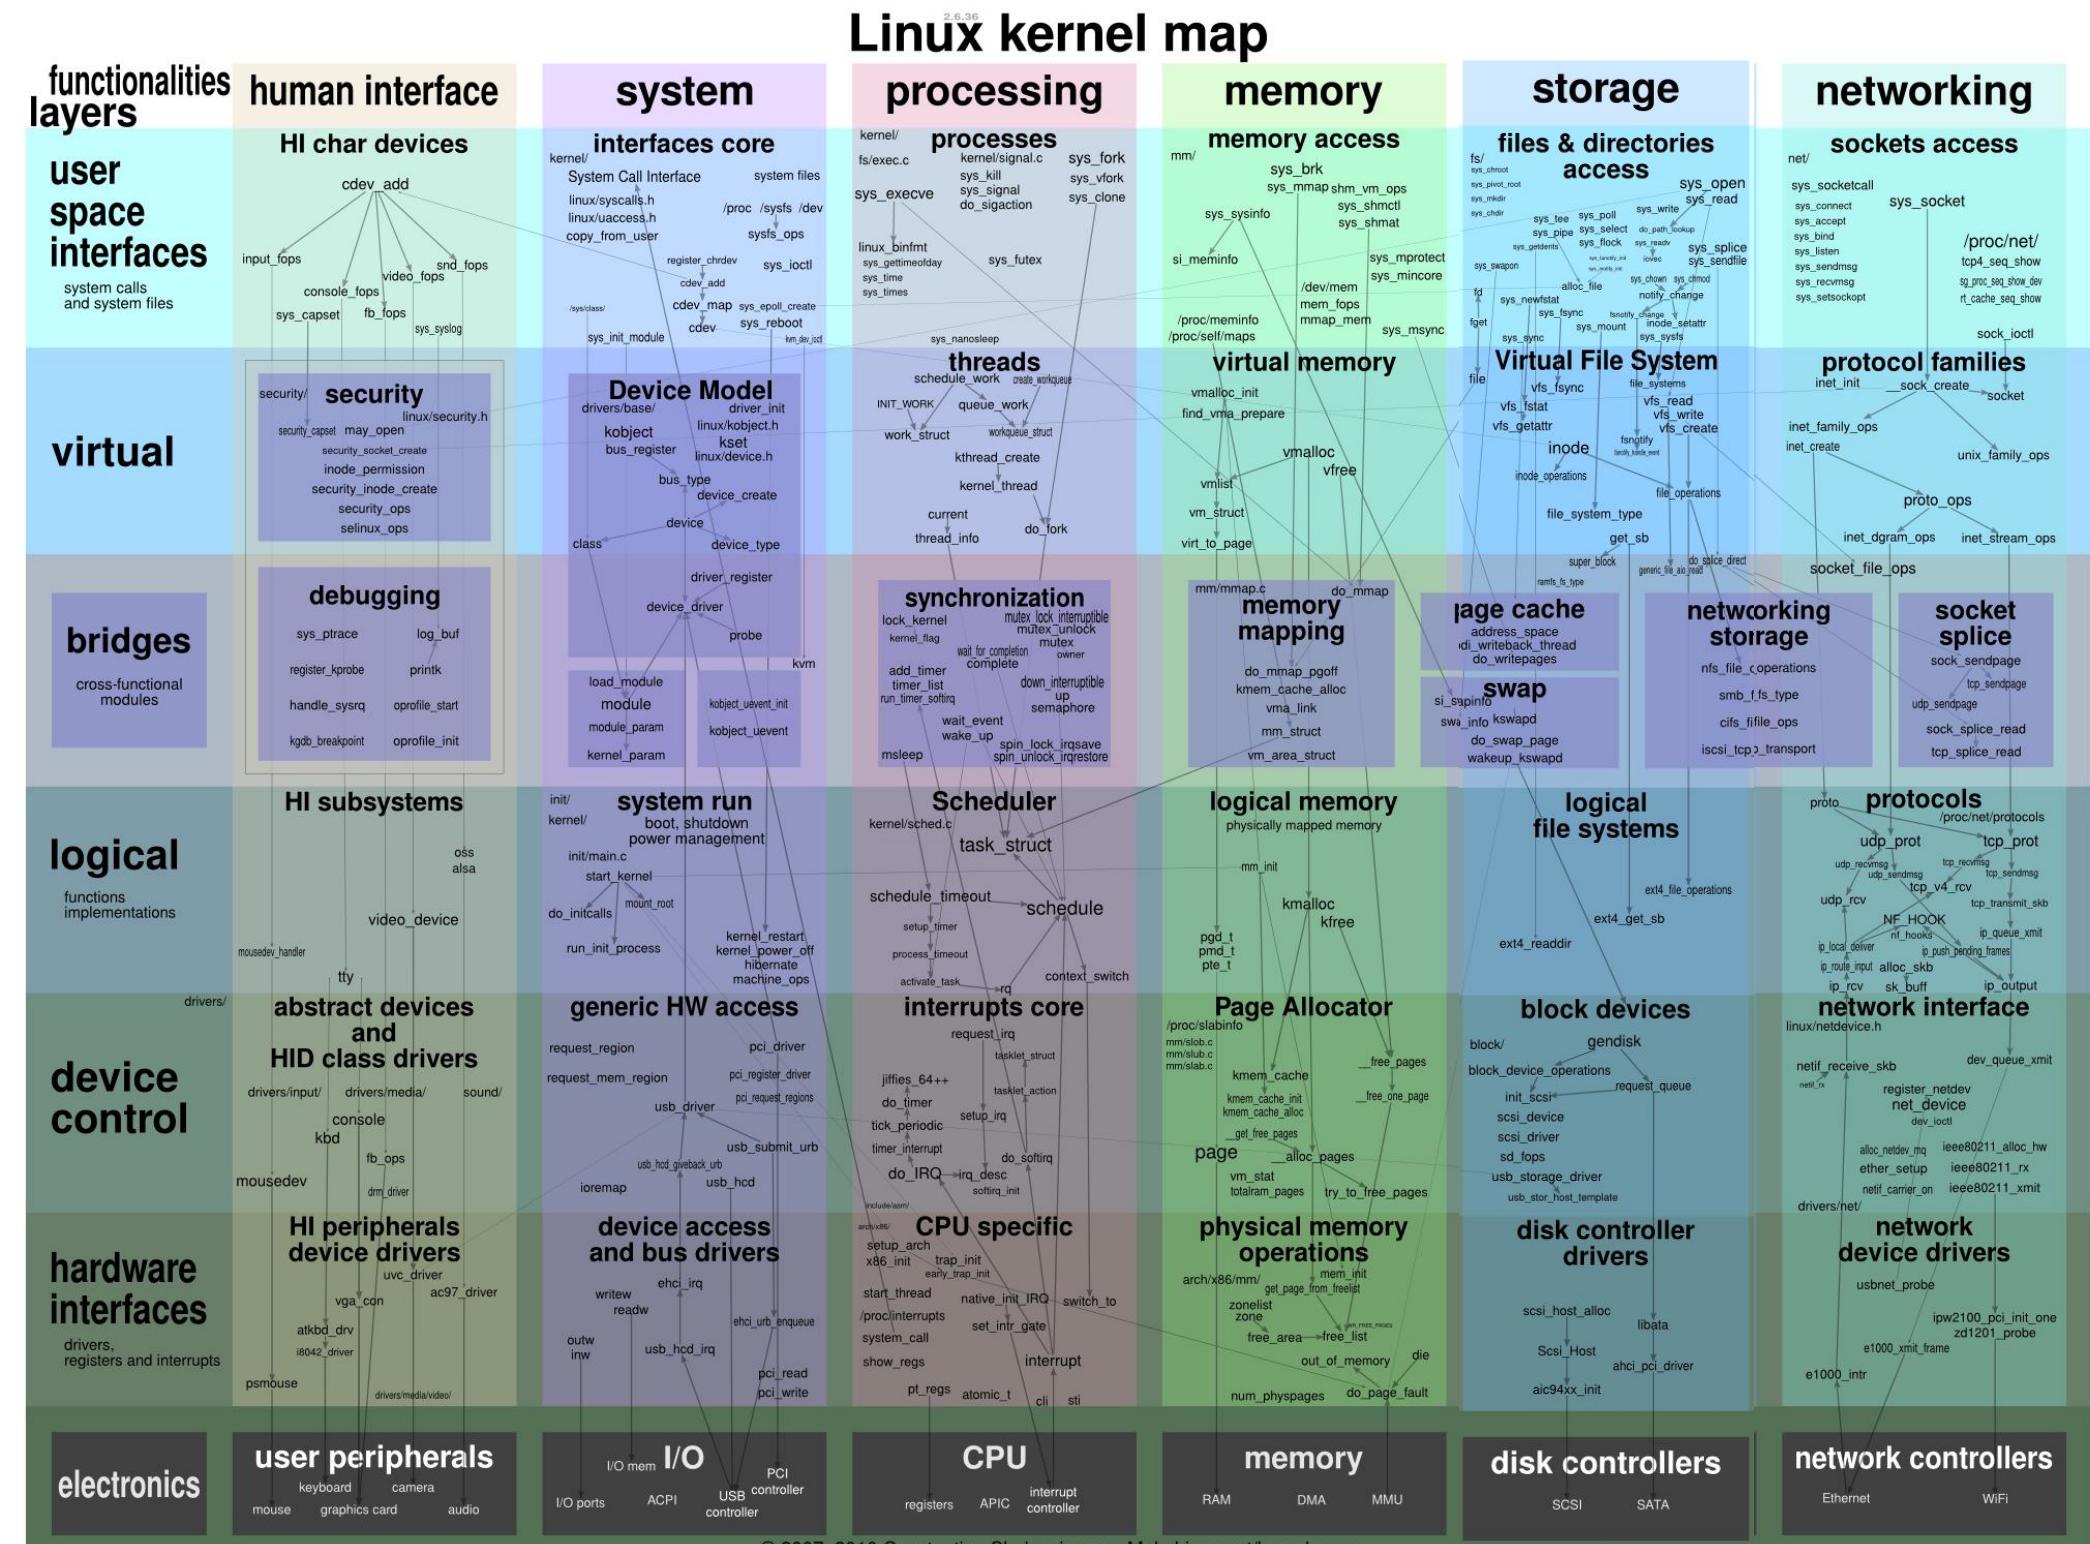
- Processes must have memory to play with
- Physical memory must be shared
 - processes must be protected from each other
 - the system must be protected from processes
 - processes must be protected from themselves?
- Process address spaces can be significantly larger than what is physically available!
 - what now?



Cache, cache, cache...

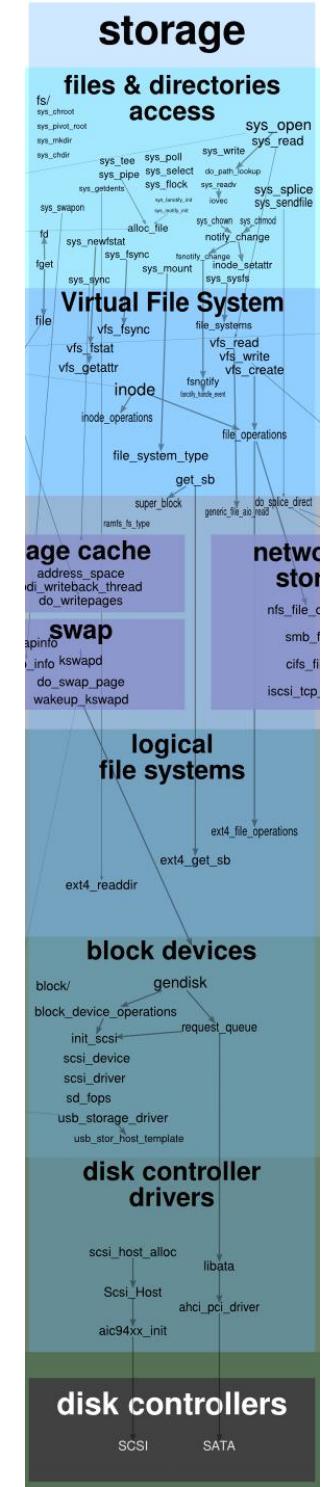
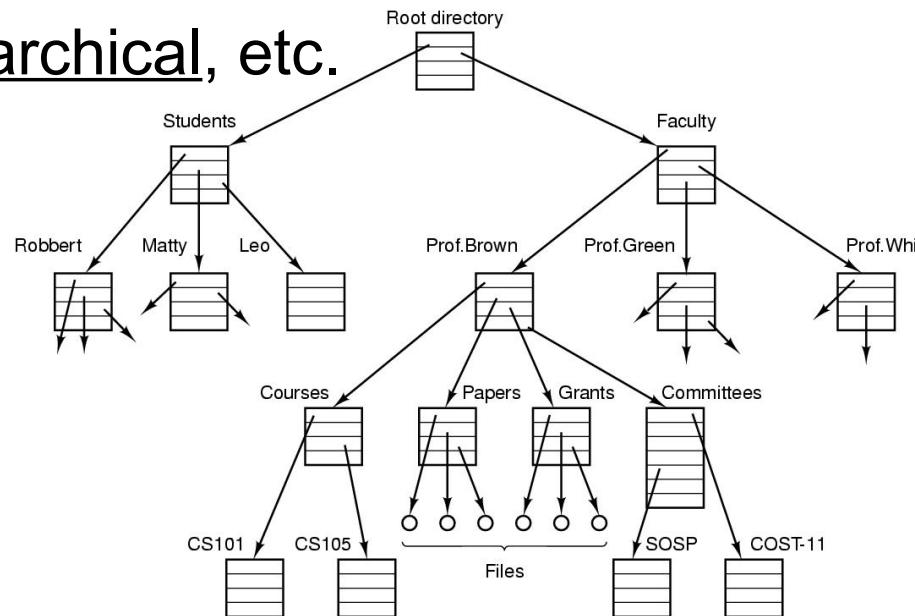
- Registers are fast, expensive
- Memory access is slow (by comparison)
- Caching is critical
 - if the cache contains the needed word: 2 cycles
 - if not: many many cycles
- So?
 - When to cache?
 - What should be cached?
 - How is stuff evicted from the cache?
 - Where does evicted stuff go?
 - How can we make this efficient/fair/safe?





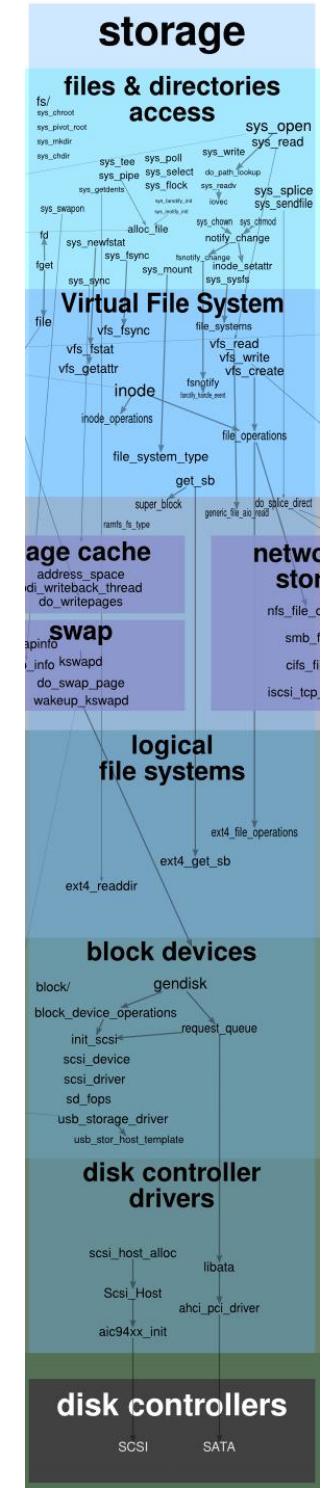
File Systems

- How is data organized?
- Files, directories, how to define these in the OS, and how to interact with them?
 - associated system calls, structures, etc.
- Organizations
 - flat, hierarchical, etc.



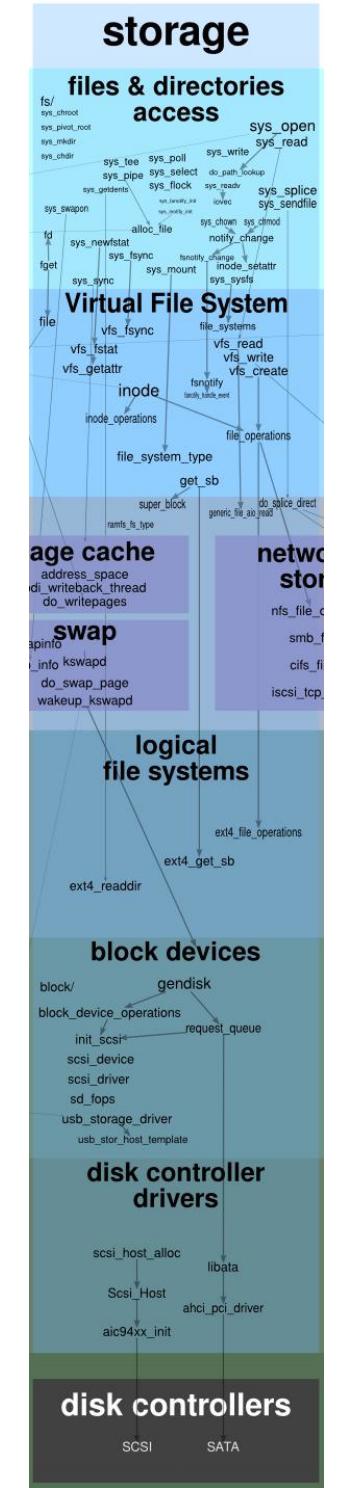
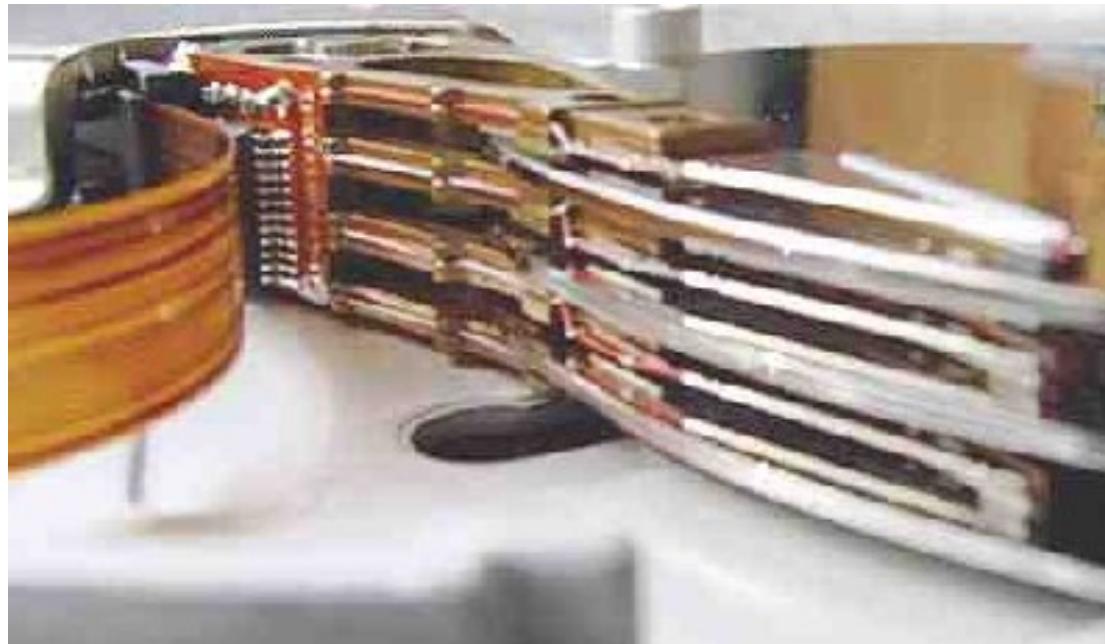
File Systems

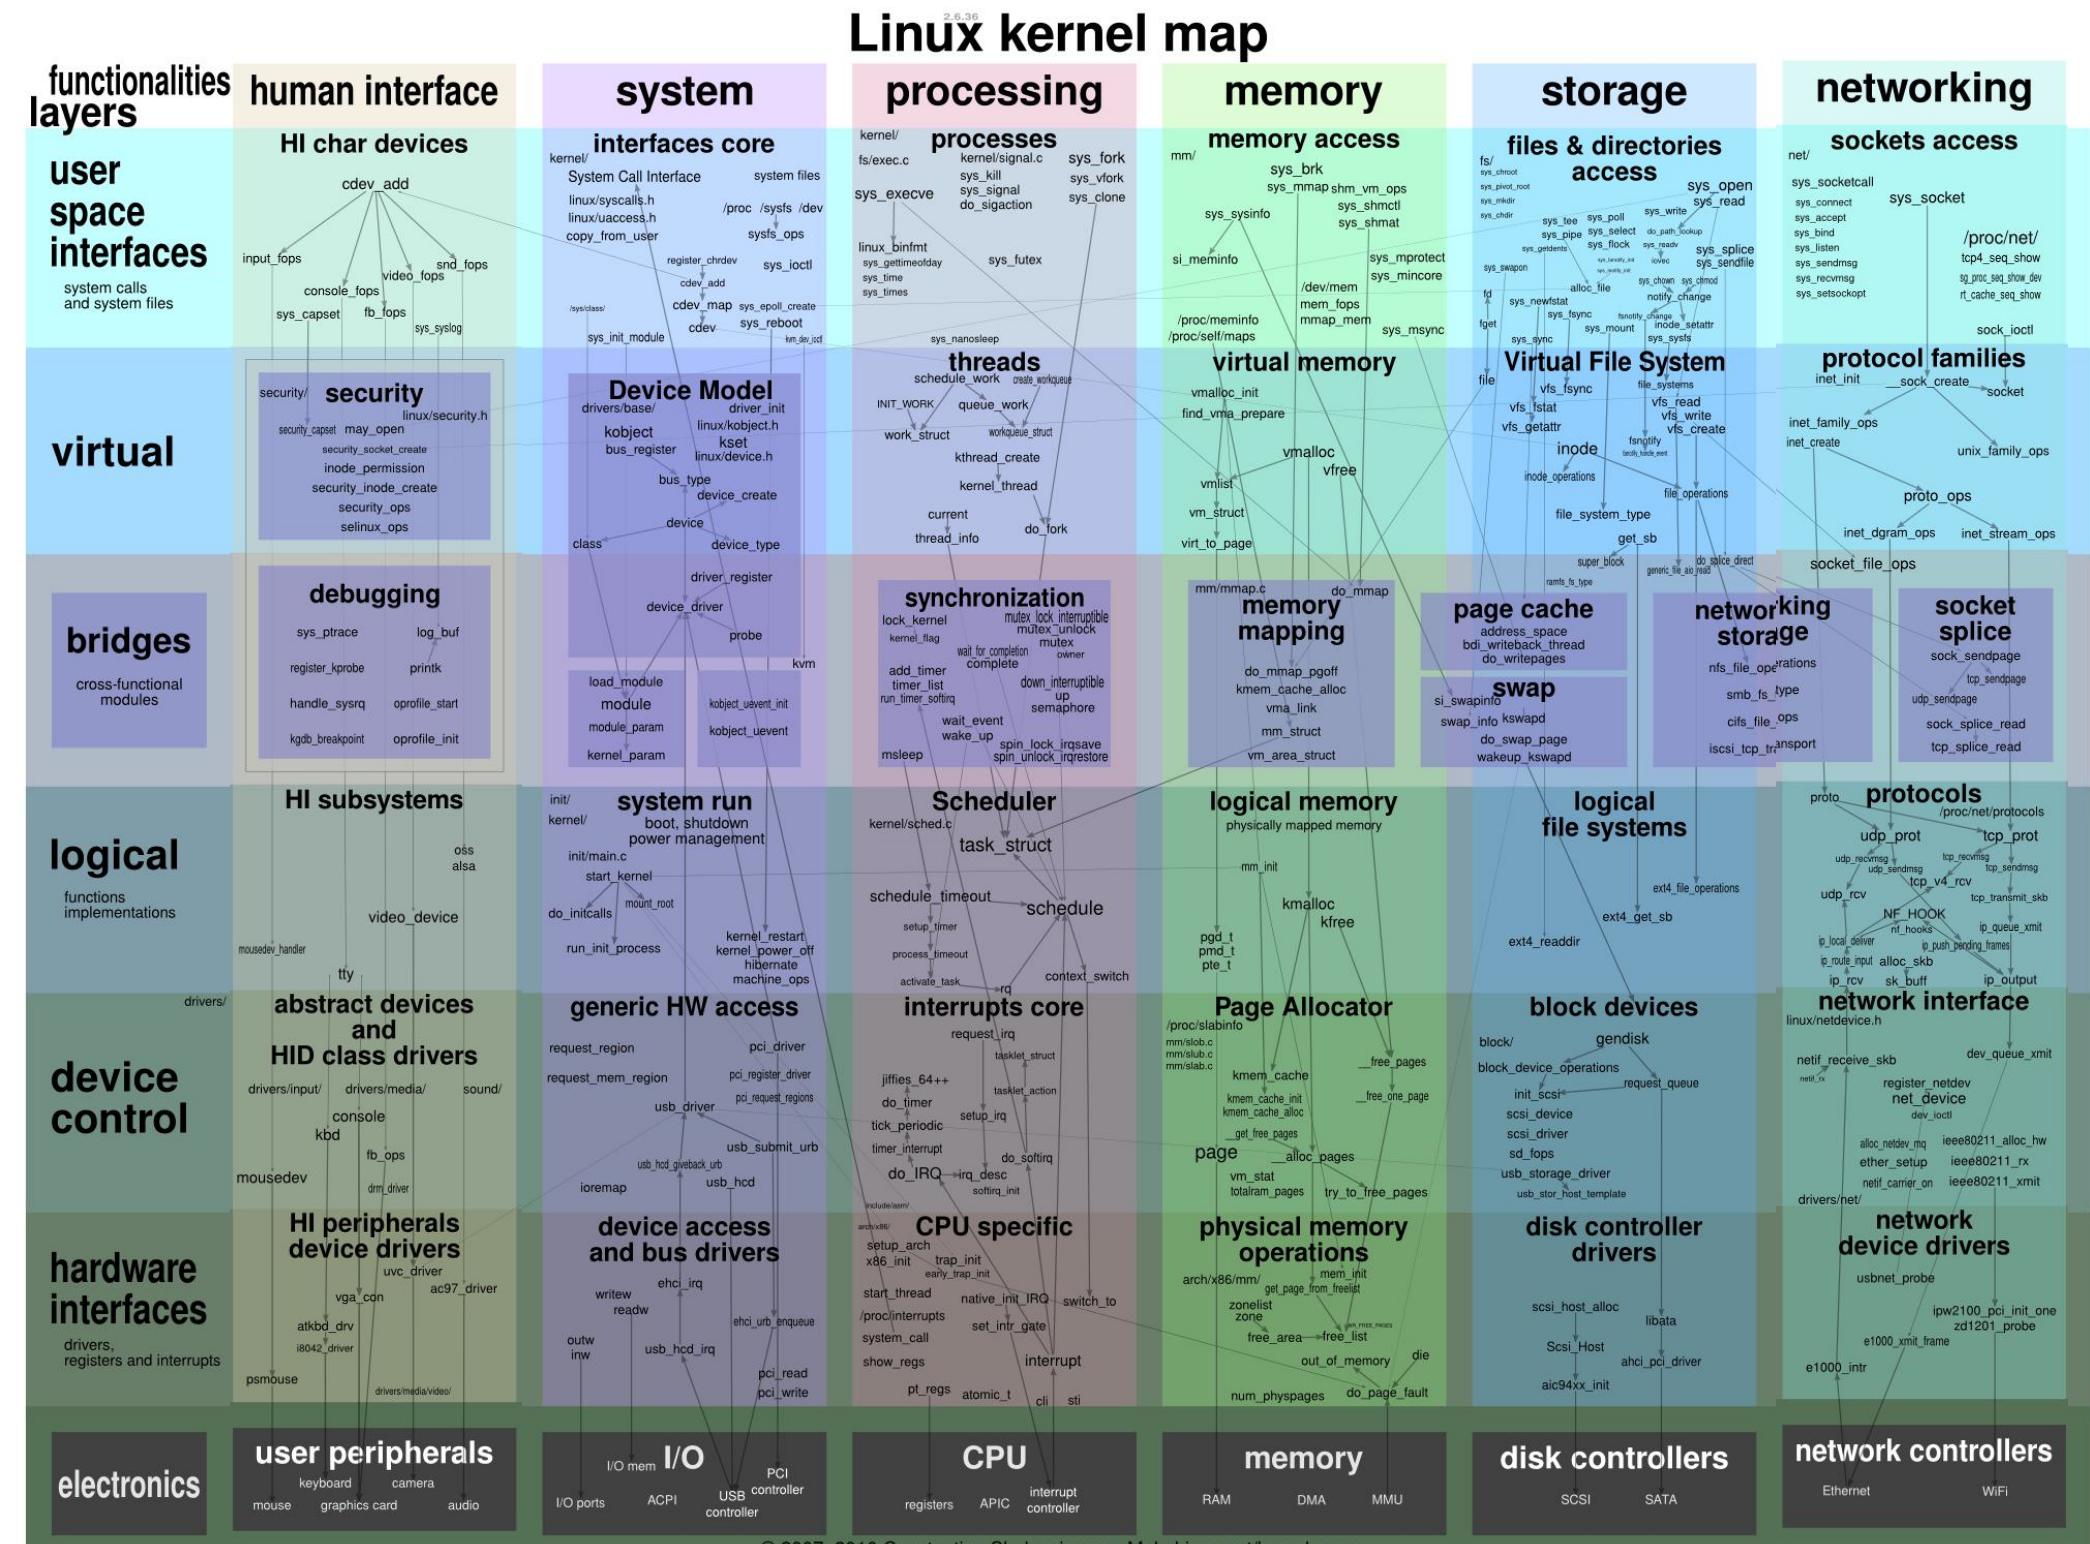
- Representation of files/directories in the OS
 - inodes
 - design choices directly affect performance
 - System call interfaces for working with files
 - Fun with hierarchical file systems
 - hard/soft links, mounting, pipes, etc.



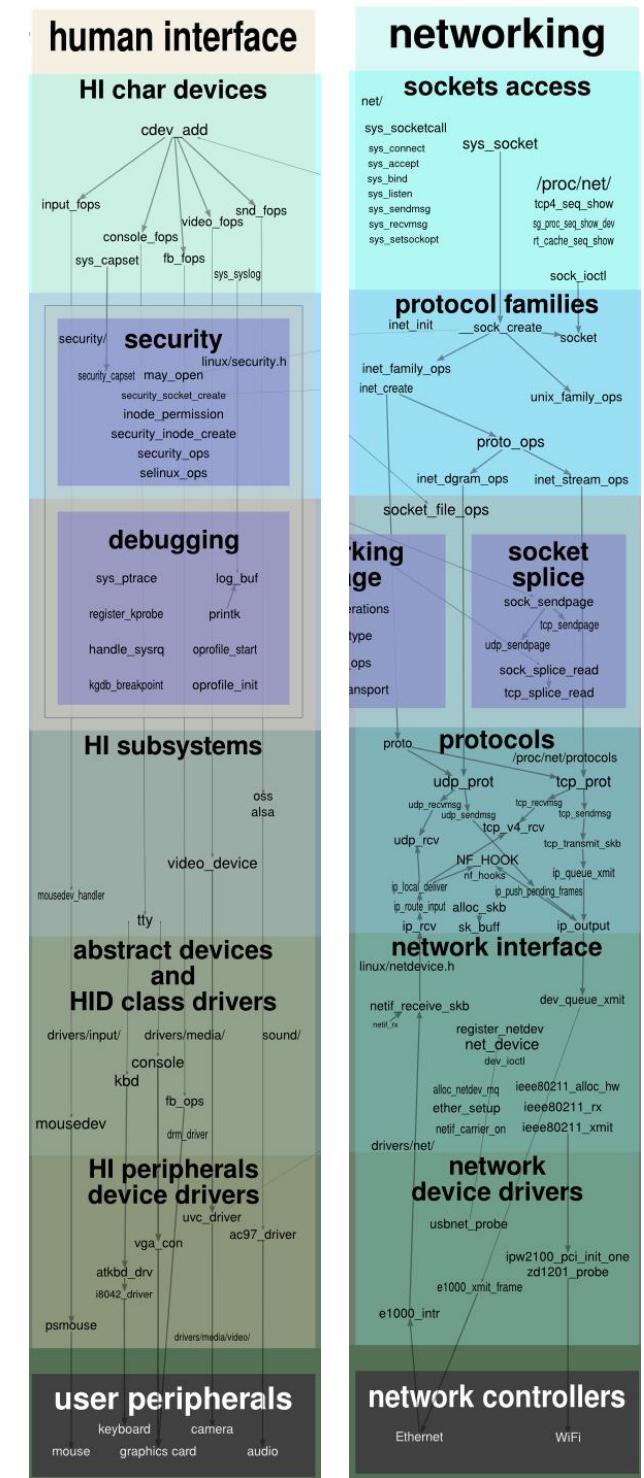
Disks

- Data organization, optimization, performance
 - take advantage of the physical properties
 - Effectively extend physical memory
 - Need caches to hide their awful slowness





- As time allows...



Other Topics

- Energy considerations/management
 - how to extend battery life, reduce cost
- Security
 - how systems get compromised
 - how can systems be protected
- Android case study
- Windows case study

Resources

- Class
 - <https://classes.cs.siue.edu/>
 - <http://www.cs.siue.edu/~icrk>
- Linux/bash
 - <http://www.makelinux.com/home/>
 - www.makelinux.com/books/Bash-Beginners-Guide/
 - <http://www.makelinux.com/books/abs-guide/>
- machine
 - os.cs.siue.edu
 - From unix-like systems, open a terminal and type
ssh <your SIUE username>@os.cs.siue.edu
 - From windows, use PuTTY, cygwin, or ...