

1 About Project 3

You have a choice. Option 1: use what you have learned about DLXOS through the past 3 projects to implement memory paging with single and two-level page tables. Option 2: use what you have learned about file I/O over the past several homeworks to implement a file system inside of a file.

The two options are conceptually very similar. Option 1 requires significantly less code to complete, is to be done in DLXOS, and is tricky. Option 2 is to be done entirely in C, as a regular Linux program, it demands more code, but you can fully enjoy the benefits of gdb. It is also tricky.

Due to the shortened timeframe and no possibility of an extension, 80 points will constitute 100% for this project's grade. Any additional points beyond 80 are extra credit. For even more extra credit, complete both options.

2 Option 1

2.1 Overview

We will be extending DLXOS's memory management subsystem to support dynamic one-level paging and dynamic two-level paging, so there are 2 parts to this project.

DLXOS has been changed to fit the requirements for this lab. You will have to copy this new version of DLXOS (`~/Public/cs314/project3.tgz`) to see the changes.

2.1.1 Setup Instructions

As before, but unpacking `project3.tgz` will give you two directories this time: `part1` and `part2`. There are 2 parts to this project. Place your solution to each part in the respective directory.

Make sure you have `~/Public/cs314/dlxtools/bin/` in your PATH.
Compile as before: type `make` in `part1/src` or `part2/src` directory.
Run in `part1/execs` or `part2/execs` directory.

2.2 Assignment

DLXOS currently supports a static one-level page table, which allocates a single 64KB page (for all of the code, data and the stack segment) per process. This page is allocated when the process is first created. The total amount of physical memory present in DLXSIM is 2 MB.

2.2.1 Part 1: Dynamic One-Level Paging

Extend the current implementation to support dynamic one-level page tables. The system should support 256 pages of 8 KB each in total. The virtual address space of each process is limited to 512 KB. During process creation only the necessary physical pages (3 pages for text and data, and one page each for user stack and system stack) are allocated. The rest of the physical pages are dynamically allocated when a page fault occurs. Once a physical page is allocated, it is not freed. A process can use at most 16 physical pages during its execution. The kernel kills a process if it exceeds the 16 page physical memory limit and displays an informative error message.

Modify the constants `MEMORY_L1_PAGE_SIZE_BITS` and `MEMORY_L2_PAGE_SIZE_BITS` appropriately.

Your implementation should be in the `part1` directory. You may test your program by running the test cases given in this directory.

See the notes below for some useful details.

2.2.2 Part 2: Dynamic Two-Level Paging

Extend the one-level dynamic page table implementation to support a two-level dynamic page table. The virtual address space of each process is limited to 16 MB. Any virtual address referenced by a process is translated to a physical address using the two-level page table. Each L1 (Level-1) page table entry refers to a L2 (Level-2) page table and each L2 page table entry points to an 8 KB page. Each process can use at the most 1 MB of physical memory during its execution (this excludes the pages being used for L2 page tables). The kernel kills a process if it exceeds the 1MB physical memory limit and displays an informative error message.

	L1		L2		OFFSET	

Each L1 page table entry corresponds to 512 KB of the total addressable virtual memory. There are a total of 32 entries in the L1 table (16 MB / 512 KB). The L1 part of the address is therefore 5 bits.

From the above we know that there are 32 L2 tables. Each L2 page table entry corresponds to 1 page (8KB) of the total addressable virtual memory. Hence, there are 64 entries in each L2 table ($16 \text{ MB} / (32 * 8 \text{ KB})$) and each L2 part of the address is 6 bits.

The offset is designated by the remaining 13 bits.

In DLX `MEMORY_L1_PAGE_SIZE_BITS` is amount mapped by a single entry in the L1 table and `MEMORY_L2_PAGE_SIZE_BITS` is amount mapped by a single entry in the L2 table. Hence

```
#define MEMORY_L1_PAGE_SIZE_BITS 19 // each entry in L1 represents 512 KB
#define MEMORY_L2_PAGE_SIZE_BITS 13 // each entry in L2 is 8kb
```

If a particular entry in the L1 table does not exist, that entry should be zero. L2 PTEs point to real physical pages in memory. In addition to the physical page address, there are three special bits in each PTE: valid, dirty, and reference bits. The valid bit must be set by the operating system to indicate that the page in memory is valid. The dirty and reference bits are set by the "hardware" when an address in the page is modified and accessed (read or written), respectively. The bits are never reset by the "hardware", but they may be reset by the OS if desired.

In order to activate two-level page tables, appropriately modify the constants `MEMORY_L1_PAGE_SIZE_BITS` and `MEMORY_L2_PAGE_SIZE_BITS`.

To simulate 16MB of physical memory, you need to add the flag `'-m 16777216'` as shown below:
`dlxsim -m 16777216 -x os.dlx.obj -a -u userprog.dlx.obj`.

See the notes below for some useful details.

Do your implementation in `part2` directory. You may test your program by running the test cases given in this directory.

2.2.3 Notes

- When a process is created it will need 3 pages for its code and data segment, one page for the user stack, and one page for the system stack. The user stack should always be initialized to the last page in the virtual address space of a process.
- The `userprog1.c` that is provided, is supposed to introduce 12 page faults. So, with initial 5 pages, this should end the program with an error(given that a process can have only 16 pages at a time). But do note that the page allocated for system stack is not counted in above mentioned math. Only the pages for code/data segments and user stack are referenced through the page table. As in the given code, you can see that the page allocated for system stack is not associated with the process's page table at all, thus making it not countable in the 16 allotted pages. So, 12 page faults would not result in an error but a 13 would as in `userprog2.c`.

- The highest virtual address space for a process is defined as `TOP_VIRTUAL_ADDRESS_SPACE` in `memory.h`. The user stack pointer initialization is dependent upon it. You can find the dependency by looking at `ProcessFork()`.
- A `TRAP_PAGEFAULT` trap is generated when a process tries to access an invalid address. A "page fault handler" `PageFaultHandler()` is declared and called in the case of `TRAP_PAGEFAULT`. Implement the function. After loading the missing page and fixing the page tables, return from the page fault handler.
- A `TRAP_ACCESS` trap is generated when a process tries to access beyond its virtual address space limit. `ProcessKill()` is called to handle the trap in `traps.c`. Implement the function so that only the process is killed and the simulator does not exit.
- When a process exits, all physical pages allocated to a process need to be freed (in case of two-level paging free the pages used for L2 page tables also). `ProcessKill()` is called in the trap of `TRAP_EXIT` to clean the process.
- You may look at the function `CPU:VaddrToPaddr` in `dlxsim.cc` to see how virtual address to physical address translation is done.
- If the pagefault handler successfully allocates a physical page, make sure to print the page faulting address inside the handler. This should be of the following form: "Process id ..., page fault address: 0x....., physical page number allocated:".
- When a process is killed or when a process exits, you should print all the physical page numbers being freed. This should be of the following form: "process id ..., virtual page base address: 0x....., Freeing physical page number: ..., "
- All virtual page addresses should be displayed in `%x` (hexadecimal) format. All physical page numbers should be displayed in `%d` (decimal) format.
- `MemoryTranslateUserToSystem()`: This procedure is used to translate user addresses into system addresses. The current implementation supports only static one-level page tables. Modify it to handle both one-level and two-level page tables.
- `ProcessKill(PCB * pcb)`: This system call should release all the resources (i.e. all physical pages) held by the process and then remove its context. You will need to use the function `ProcessFreeResources()` already defined in `process.c`. You will need to change the implementation of `ProcessFreeResources()` for both parts of this project appropriately.

2.3 FAQ

Q: Debugging is a nightmare, how can I tell what is going on?

A: Make good use of the debugging `printfs`, they can be enabled individually by using the tag associated with each `dbprint`. You can also enable all of them with `'+'`. Should look something like this: `dlxsim -x os.dlx.obj -a -D + -u userprog.dlx.obj`

Q: When the userprog.dlx.obj starts it displays the usage message: `Usage: [case id]` then it is followed by repetitive error lines: `Got an open with parameters (' [case id] ',0x0)`

It seems that userprog isn't receiving the arguments when I run it.

A: Don't mess with the sysStackArea. The code that allocates the page for sysStackArea should not be touched. Also, the "Got an Open" messages usually come from when dlxos tries to run a process that is done or not runnable. Have you implemented the ProcessKill function yet? That is what is called when a process exits now instead of ProcessDestroy, but it currently does not do anything. Hope one of those prompts points you in the right direction.

Q: In the assignment it says that ProcessKill() is implemented so that the simulator does not exit. Does this include the case when there are no more processes to run?

A: Call ProcessSchedule() at the end of ProcessKill(). If there is no runnable process, let the system exit. Print whatever reasonable messages you think. Make sure to give yourself, and me, a clue what's going on from the message.

Q: In the slides it says that we need to modify MemoryTranslateUserToSystem(), but in the code the comments say "This works for simple one-level page tables, but will have to be modified for two-level page tables." Does this mean that we do not have to modify MemoryTranslateUserToSystem()?

A: "simple one-level page table" means it works for dlxos with a single 64kb page. You need to think about what to change when moving to multiple pages.

Q: What does MemorySetFreemap () do?

A: It's called to flag a physical page is free. You shouldn't call it directly. Instead, use MemoryFreePage().

Q: In memory.h, do we need to change the memory_max_pages?

A: You don't change memory_max_pages. That's for physical space.

Q: I am getting a page fault at location 0. This is very weird. Is this normal?

A: Most likely you didn't set table entry properly. Remember to use MemorySetupPte().

Q: What value should TOP_VIRTUAL_ADDRESS_SPACE be?

A: It determines the top of virtual address. Valid virtual address should be smaller than it. For part 1 it should be: `#define TOP_VIRTUAL_ADDRESS_SPACE 0x40 * MEMORY_PAGE_SIZE`
For part 2 it should be:
`#define TOP_VIRTUAL_ADDRESS_SPACE 0x800*MEMORY_PAGE_SIZE`

Q: In addition to the physical page address, there are three special bits in each PTE: valid, dirty, and reference bits. The valid bit must be set by the operating system to indicate that the page

in memory is valid. The dirty and reference bits are set by the "hardware" when an address in the page is modified and accessed (read or written), respectively. The bits are never reset by the "hardware", but they may be reset by the OS if desired. Its not clear to me which bits we will be setting/changing and which bits will be set by the system. Also, it seems to me that if the PTE will keep track of these bits, then we need to know the format of where these bits will be in the physical address (ie at the end, beginning) I imagine that the address will be shifted to the left 3 bits and this area will store the three above bits, but since the "hardware" will be using these bits as well I'd like to know where/how they're expected to be implemented.

A: The bit you need to handle is valid bit. You should set it when allocating a page, reset it after freeing the page. The mask for valid bit is `MEMORY_PTE_VALID` in `memory.h`

Q: Does the `MEMORY_MAX_PAGES` value need to change in `lab4_2`? It seems like we would since we're addressing 16MB instead of 2MB.

A: Don't worry about it. We won't use that much space.

Q: I'd like to know more about L1 and L2 in part 2.

A: In part 2, you should only store L1 table in PCB and allocate page for L2 table dynamically when needed. Here's an outline:

1. Allocate a page for L2 page table.
2. Setup L1 entry for L2 page table using `MemorySetupPte()`
3. Calculate L2 page address: `L2 = L1[L1_page_no] & (MEMORY_PAGE_MASK)`
4. Allocate pages for actual text, data, or stack.
5. Use L2 page address to get pointer to L2 table and setup L2 entries for text, data, or stack pages using `MemorySetupPte()`. e.g. `L2[3] =`

2.4 Grading Rubric

20 points for reasonable attempt

Part 1:

```
dlxsim -x os.dlx.obj -a -u userprog1.dlx.obj
```

20 points possible

```
dlxsim -x os.dlx.obj -a -u userprog2.dlx.obj
```

20 points possible

Part 2:

```
dlxsim -x os.dlx.obj -a -u userprog1.dlx.obj
```

20 points possible

```
dlxsim -x os.dlx.obj -a -u userprog2.dlx.obj
```

20 points possible

In each of the above, see the userprog source for a comment detailing how the program should behave. If the number of page faults is off, or the program does not exit/exits when it is not supposed to (sometimes the program should be terminated due to a lack of pages), or if the program fails to run, will all be cause for point deductions.

3 Option 2

3.1 Overview

The Old File System (the original file system developed at Bell Labs) is the granddaddy of Unix filesystems. It divides each drive into a number of partitions, each of which may contain a file system. One file system never spans multiple partitions. A file system is described by its superblock, which contains the basic parameters of the file system. These include the number of data blocks in the system, a count of the maximum number of files, and a pointer to the free list, a list of the free blocks in the file system.

Within the file system are files. Certain files are distinguished as directories and contain pointers to files that may themselves be directories. Every file has a descriptor associated with it called an inode. An inode contains information describing the file and an array of indices pointing to data blocks for the file. A file's contents can be produced by specifying a file name, which is resolved to the specified file's inode by parsing the individual parts of the file's absolute path and retrieving directory data at each step as needed.

Some basic command line argument parsing is included in the project handout file at `~/Public/cs314/project3.tgz`.


3.2 Assignment

Implement a restricted version of the Old File System as described in the sections that follow. You are to implement it in C and are allowed to use any standard libraries. The file system is to reside not in a partition on a disk, but rather in a single fixed-size file. The interface to your file system will be a program with command-line options for writing, reading, removing, and listing files contained in the file system.

3.3 Formatting

The fixed-size file that is acting as a stand-in for a disk partition will need to be formatted to provide some basic information about your file system each time it is opened by your program. The formatting will result in a superblock, free list, and inodes being written to your fixed-sized file. The superblock defines the file system, specifying the number of blocks available in the file system, the size and number of inodes, and potentially other useful information. The most straightforward free list might be implemented as a bit map of sufficient size to provide a bit for each block in the file system, specifying whether the corresponding block is free or used. Inodes should initially be unused. The blocks that are used for file data storage will appear after the superblock, free list, and inodes. No action is required for the blocks themselves, since the only way they might be read or written is when the file system structures like the free list or inodes refer to them. Make the size of your file system 10MB. Do not turn your file system file in with your solution.

You can think of the file your file system will reside in as having roughly this structure:



3.4 File Types

Your file system should support two file types, directories and regular files. Directories are basically implemented as regular files that are labeled as directories and only contain directory entries. Don't worry about compacting directories as entries are added and removed, instead think about marking directory entries that belong to removed files as invalid, and potentially reusing them when a new file is written to that directory. No format should be assumed for regular files, you are only expected to be able to read and write these in their entirety. Keep in mind that files (and the files we call directories) will commonly span more than one block, and that the blocks need not be contiguous.

3.5 Inodes

Your inodes are to contain 100 direct block references. Don't worry about indirect blocks. You are not reproducing a user environment, so you also don't have to store the user, group and permissions information. I also don't care about time stamps. Just focus on what will allow you to read and write a file in its entirety. Each block that the inode points to should be 512 bytes (your file system will apparently only handle smallish files). I will not be testing with more than 100 files, so have at least 100 inodes in your file system.

3.6 Functionality

Your program should allow files to be added to your file system, read and removed from it, and for the contents of the file system to be listed. It might be easiest to describe the intended functionality by providing examples of each. For example, the following command line invocation should list the contents of the example file system starting with the root directory:

```
./filefs -l -f example
```

The output might be a file system tree, showing the names of each directory and file contained within your file system (think about recursing through it starting with the root, adding a number of leading tabs to each line of output equivalent to the depth of recursion).

The following invocation would add the specified file to your file system (as well as each directory in the path):

```
./filefs -a /a/b/c -f example
```

In this example, `/a/b/c` is the absolute path to an existing regular file `c` in `/a/b/`. The intended result is that your file system will likewise contain directory `a`, subdirectory `b`, and file `c`, with `c`'s contents being identical to the specified file's. So, you'd find some free inodes and blocks and write directory entries to both `/a` and `/a/b`, then get another free inode for `c`, and copy `c` into your file system, repurposing free blocks and updating the block mapping in its inode as you perform the copy.

The following invocation would remove the specified file from your file system (as well as any directory that becomes empty after you remove the specified file):

```
./filefs -r /a/b/c -f example
```

So, if the only file in your file system is `/a/b/c` and you remove it as above, then listing the contents of the file system should result in nothing being printed, since both `a` and `a/b` would have been removed after `c` was removed.

The following invocation would read the contents of the specified file to stdout, then redirect that output to file `foobar`.

```
./filefs -e /a/b/c -f example > foobar
```

Reading a file's contents in this manner should not result in the file being removed from your file system. There should be no difference between the original file and the data returned by your file system.

3.7 Hard Requirements

- Your file system should be self-contained; i.e. metadata and data comprising the file system should be inside of a single file, formatted according to your interpretation of the specifications.
- Your file system should work for any kind of file, text or binary.
- Your file system should be able to support paths or filenames up to 255 characters total.
- Do make sure that if you store a file in your file system and then extract it, it matches the original file (perform a diff on the original and extracted file to make sure there are no differences).
- Include an up-to-date Makefile. I don't want to have to reproduce your build.
- Make sure it compiles before you turn it in.

3.8 Grading

- 20 points if there is a reasonable attempt at a solution conforming to the original problem statement.
- 20 points if the file that is created to serve as the storage for your file system is formatted correctly.
- 15 points if files can be added to the file system (directories should be created in your file system as the specified input file's path indicates).
- 15 points if the contents (directory and file names) of the file system are displayed when a listing is requested.
- 15 points if files can be removed from the file system. (listing contents of the file system no longer shows removed files)
- 15 points if files can be extracted from the file system. (file data returned via stdout matches the file data of the original that was previously added to your file system)

4 What to Turn In

Turn in 3 things as a group:

- group.txt containing information about your group.
- A .tgz of your project directory, as before. (Include a design.txt file in which you explain where you made your changes and what the changes are intended to do (option 1) or what your design looks like and how you implemented it (option 2). Any additional information that would be useful for grading: testing, configuration, compiling, etc should go into a README file.)

Individually, you should turn in a text file containing a brief account of your group activity. Keeping in mind your group's dynamic, answer at least the following questions: What worked? What did not work? What were you responsible for in the project? What could be done differently next time?